

Crowding

Yizhen Gu

Qu Tang

Kai Wu

Ben Zou*

September 2026

[Click here for the latest version]

Abstract

Crowding is a common disamenity, yet little is known about its monetary cost. We estimate the willingness to pay (WTP) to avoid crowding in public transit with a revealed-preference approach. We use departure-time choices in the Beijing Subway, where an early-bird discount generates fare variation, to estimate how passengers trade off fares, in-train crowding, and schedule deviation. Leveraging the subway's network structure, we impute real-time in-train crowding, infer each passenger's optimal arrival time, and construct a novel instrumental variable for crowding. The marginal WTP to reduce crowding by one passenger per square meter is about 0.05 RMB per minute, implying a crowding externality comparable to the fare. High-income passengers have higher WTP to avoid crowding but are less price-sensitive. An optimal crowding tax improves welfare within the subway system but is regressive; a two-class configuration elicits self-selection and benefits both high- and low-income commuters.

Keywords: Crowding externality, Public transit, Travel demand.

JEL Codes: R41, R48, L92, D62

1 Introduction

Many goods decline in quality as more people use them at the same time. Consider a grocery store on a Saturday afternoon or a beach on the first hot day of summer. As more people arrive, individuals

*Gu: Peking University, yizhengu@phbs.pku.edu.cn; Tang: Jinan University, qutang@jnu.edu.cn; Wu: University of Texas at Austin, kaiwu@utexas.edu; Zou: Purdue University, zou136@purdue.edu. We thank Marika Cabral, Yue Cao, Jackson Dorsey, Gilles Duranton, Matthew Freedman, Ryuya Ko, Shanjun Li, Daniel McMillen, Hayden Parsley, Cody Tuttle, Qinshu Xue, and workshop and conference participants at the Chinese University of Hong Kong (Shenzhen), Shanghai University of Finance and Economics, the University of Texas at Austin, the Asian Bureau of Finance and Economic Research Conference, and the Urban Economics Association Annual Meeting for comments. Yunlei Tang, Danhe Tang, and Xiangjie Zhu provided excellent research assistance. Gu gratefully acknowledges funding from the Natural Science Foundation of China (Grant No. 72273008). All remaining errors are our own.

have less room and feel increasingly uncomfortable. This is crowding, defined as sharing a limited space with many others. Each additional person makes the space more crowded and increases the discomfort experienced by others, imposing a negative externality.¹ Quantifying this externality requires knowing how much people are willing to pay to avoid crowding. We estimate this parameter for crowded public transit during the morning peak, using the Beijing Subway as our setting.

Crowding on public transit is common in densely populated areas and can be severe during peak hours. Improvements in commuting conditions, including reduced crowding, can improve worker performance (Lu et al., 2024). Dissatisfaction with crowding is also associated with greater stress and lower life satisfaction (Kumagai et al., 2021). Crowding also shapes how transit services are provided. In addition to expanding capacity through more frequent services or larger vehicles, operators can use fares to encourage travel outside peak hours, offer different service classes, or even limit passenger entry through queuing when crowding becomes severe.² Estimates of marginal willingness to pay (MWTP) to avoid crowding are a key input for evaluating these policies.

Yet credible estimates of the willingness to pay for reduced crowding alone remain limited.³ There are at least three empirical challenges. First, detailed data on the crowding individuals experience are rarely available. Crowdedness changes as people enter and leave, so measuring exposure requires knowing how densely occupied a space is and how long each person spends in it. Most existing studies rely on stated-preference surveys rather than observed choices. Second, crowding often varies with other service attributes, making its separate value difficult to identify. Buses are more likely to be crowded when roads are congested, leading to delays in travel. Less crowded options, such as business class on airplanes and skyboxes in stadiums, typically come with other perks, including service quality and status that are difficult to measure. Observed choices therefore do not reveal how much people are willing to pay for reduced crowding alone. Third, crowding is endogenous. Crowding arises when many people respond to the same, often unobserved, demand shock and converge on the same location at the same time. An exciting game or a holiday sale can attract more people while also making attendance or shopping more valuable. Crowding may therefore be positively associated with demand even when people dislike it.

We address these challenges using administrative smartcard trip records for the Beijing Subway during the morning peak in 2016.⁴ These records report each passenger’s entry and exit stations and

¹In some contexts, crowding may even be perceived as an amenity, as with the “vibes” at busy bars and music concerts.

²For example, London’s Piccadilly line upgrade includes plans for more frequent peak service. Hong Kong’s MTR offers an early-bird fare discount to encourage passengers to travel earlier. In the Tokyo region, JR East offers premium Green Cars alongside ordinary cars on commuter trains, with more spacious seating available for an additional fare. In Shenzhen, busy metro stations restrict entry during peak hours, requiring passengers to queue before entering.

³Much of the crowding-valuation literature relies on stated-preference surveys (Li and Hensher, 2011). Revealed-preference studies use observed route choices to estimate trade-offs between crowding and travel time (Hörcher et al., 2017; Yap et al., 2020). These studies provide evidence on crowding valuations, but leave limited evidence on monetary willingness to pay identified through observed fare variation.

⁴Smartcard payments accounted for an estimated 90–95 percent of subway trips around the time of the EBD.

times but not where they were in between. We develop an algorithm that combines these records with the subway network and train timetables to reconstruct passengers’ locations and measure their crowding exposure at high frequency.⁵ Subway cars are otherwise identical, much of the network was newly built, and trains have a near-perfect punctuality rate. Travel time is nearly constant conditional on the route, helping separate crowding discomfort from differences in speed and time costs, prominent confounders in transit crowding valuation. To tackle endogeneity arising from correlated demand shocks, we construct an instrument for crowding using passengers whose trips overlap incidentally with the commuters we study because of the network structure but reflect different travel demands.

Using these novel, high-frequency measures, we characterize crowding patterns during the morning peak in Beijing. Among trips originating from the 32 selected stations in our main analysis, the average density is about 3.4 persons per square meter and the ninetieth percentile reaches roughly 5 persons per square meter. The average in-train exposure lasts about 38 minutes. We show empirically that crowding does not meaningfully slow down trains or induce detour routing, supporting the interpretation of in-train crowding as nearly pure discomfort. Crowding also remains stable across dates conditional on the route and departure time, suggesting that experienced commuters can anticipate the crowdedness they will encounter.

To estimate how much commuters are willing to pay to avoid this crowding, we model how they choose when to depart, taking each commuter’s origin–destination (OD) pair and route as given. We divide departure times into 15-minute bins and estimate a discrete-choice model in which commuters trade off fares, expected crowdedness, and the cost of deviating from their optimal arrival time. Commuters can also choose an outside option consisting of other travel modes. The extent to which passengers pay a higher price for less crowding or a smaller schedule deviation reveals their MWTP to avoid crowding and rescheduling. We proxy expected crowding and travel time with fitted values from regressions on OD-by-day-of-week (DOW)-by-time-bin fixed effects. The demand model includes OD-by-date and date-by-time-bin fixed effects, which absorb daily demand differences across routes and departure-time patterns common across routes on each date.⁶ Crowding relief may be more valuable on more crowded trips and to some income groups than others, affecting both the benefits and distributional effects of policy. We examine these differences using a quadratic crowding specification and a random-coefficient model that allows sensitivity to fares, crowding, and rescheduling to vary with income.

To identify commuters’ price sensitivity, we exploit exogenous variation in fares across departure

⁵The algorithm jointly estimates each passenger’s fastest route and the average access, transfer, and exit times at each station, with in-train time taken from published train timetables. Dividing the number of passengers sharing the same route segment in the same time bin by the available train floor area yields crowding density in persons per square meter. Integrating this density over the duration of a passenger’s trip gives their total in-train crowding exposure, measured in person-minutes per square meter.

⁶These demand fixed effects allow daily conditions, such as weather, to affect demand differently across OD pairs and allow the common morning demand profile to vary by date.

times generated by an off-peak discount. The discount, known as the Early-Bird Discount (EBD), gives a 30 percent fare reduction to passengers who tap in at one of 16 stations on two suburban lines before 7 AM on workdays. Control stations without the discount help identify common departure-time patterns, allowing us to separate the fare response at EBD stations from these patterns. We provide evidence that external passengers do not systematically switch into EBD stations to claim the discount, supporting the model’s assumption that route choice is fixed.

We identify crowding disutility using an instrument that addresses the correlation between crowding and unobserved demand. The instrument is based on incidental companion passengers (ICPs), who travel from other origins to distant destinations but share train segments with a focal passenger, adding to crowding on those segments. Unobserved demand shocks change the utility of traveling on a particular OD pair, date, and departure-time bin, beyond the common rush-hour patterns absorbed by the fixed effects. We restrict ICPs to infrequent riders and assume that, conditional on the model’s fixed effects, their demand shocks are unrelated to those of the focal market, following a shared-capacity identification strategy similar to that of Hortaçsu et al. (2024) in the airline market. We define incidental companion time (ICT) as the weighted sum of ICPs, where each ICP is weighted by the duration of shared travel. The instrument is expected ICT, which we proxy using the same OD-by-DOW-by-time-bin fixed-effect projection as expected crowding. We supplement expected ICT with instruments that interact EBD-origin status with the discount cutoff and functions of departure time. Because the discount induces passengers to reschedule, it changes crowding even within periods when fares are unchanged. These instruments exploit variation in expected crowding that remains after controlling for fares and the model’s other controls.

We also need to measure how far passengers’ arrival times deviate from their optimal schedules. A common challenge in the structural travel demand literature is that optimal arrival times are unobserved and may vary both across passengers and across dates for the same passenger. We infer them from *coworkers*, defined as passengers who arrive at the same destination around the same time on the same day but live elsewhere and take routes not directly affected by the EBD. Their arrival distribution reflects arrival choices without the direct fare incentive to shift schedules. Within each destination-date cell, we reweight this distribution to match the income composition of EBD riders. We then rank EBD riders by their actual arrival time and assign each rider the coworker arrival time at the same percentile of the reweighted distribution. The method is non-parametric and uses only contemporaneous data, reducing concerns about changes in commuting patterns over time.

The baseline estimates imply substantial costs of crowding. The average passenger is willing to pay 0.053 RMB per minute to reduce crowding density by one person per square meter. Over the average in-train exposure of 38 minutes, this reduction is worth roughly 2.0 RMB, compared with an average fare of 5.3 RMB. At the average density documented above, total crowding disutility is about 6.7 RMB, or 16 percent of the median hourly wage of about 43 RMB. The MWTP rises with the level of crowding, so welfare costs are concentrated on the most congested trips. The baseline

MWTP to avoid rescheduling is 0.144 RMB per minute.

Allowing preferences to vary with income reveals substantial differences in crowding valuations. In the random-coefficient model, higher-income riders dislike crowding more and are less sensitive to fares, so their MWTP for crowding reduction rises sharply with income. At the median monthly income of 7,500 RMB, passengers are willing to pay 0.07 RMB per minute to reduce crowding density by one person per square meter; at the 90th percentile of income, 17,500 RMB, this amount rises to 1.62 RMB. As a result, few high-income passengers take the subway in the equilibrium under the current fare structure. In the random-coefficient model, the MWTP to avoid rescheduling, by contrast, is approximately 0.10 RMB per minute and relatively flat across the income distribution. High-income and low-income passengers place similar value on arriving on time, even as they differ markedly in how much they are willing to pay to avoid a crowded car.

Building on these estimates, we investigate the welfare consequences of three alternative policies: a Pigouvian crowding tax, quantity control through queuing, and a revenue-neutral two-class configuration with higher-priced Business cars and discounted Standard cars. To keep computation manageable while retaining the key policy trade-offs, we model a representative trip and abstract from network interactions and scheduling dynamics. The optimal crowding tax is 3.30 RMB, about two-thirds of the baseline fare. The tax raises welfare within the subway system, but its incidence is sharply regressive. Low-income riders lose consumer surplus, while high-income riders gain slightly even before the revenue is rebated. Subway ridership declines by 16 percent, driven by a large reduction in low-income riders. Quantity control targets the same ridership reduction but delivers lower welfare because time wasted in the queue dissipates economic rent. Both policies push passengers out of the subway, generating additional welfare losses from road congestion. Including these spillovers, welfare under the crowding tax falls slightly below baseline.

The two-class configuration, by contrast, delivers a Pareto improvement. It exploits the steep income gradient in WTP to avoid crowding: passengers who value crowding relief more pay for less-crowded service. Under revenue-neutral pricing, the optimal Business fare is 10.30 RMB, roughly twice the baseline fare, while the Standard fare is 2.44 RMB, a discount of about 52 percent. High- and low-income passengers self-select into different classes, with Standard cars more than three times as crowded as Business cars. High-income passengers benefit substantially, primarily from less-crowded cars. Low-income passengers experience far more crowded cars, but they are compensated by deeply discounted fares, and their consumer surplus increases slightly on net. The discounted Standard fare draws new riders from the outside option, and total ridership rises by about 18 percent. Attracting more people to use the subway alleviates surface road congestion, especially when high-income residents switch from cars. Including these spillovers, total welfare rises by about 15 percent.

Contributions and Related literature. This paper contributes to four strands of literature. First, we contribute to the valuation of crowding in public transit. A growing empirical literature studies demand elasticity in public transit (e.g., Davis, 2021; Gu et al., 2024; Hahn et al., 2023), but esti-

mates of within-system crowding disutility based on observed fare variation remain limited. We emphasize that the subway setting differs conceptually from the literature on road congestion (e.g., Vickrey, 1969; Arnott et al., 1993; Small and Verhoef, 2007; Anderson, 2014; Yang et al., 2020; Kreindler, 2024). Road congestion primarily generates externalities through lost time, whereas crowding in public transit generates disutility through physical discomfort even when travel times are unaffected (Tirachini et al., 2013; Wardman and Whelan, 2011; Haywood and Koning, 2015).

Much of the evidence on crowding valuation comes from passenger surveys that elicit stated preferences (e.g., Douglas and Karpouzis, 2006; Haywood and Koning, 2015; Singh et al., 2023), as reviewed by Li and Hensher (2011) and Wardman and Whelan (2011). Such surveys may suffer from hypothetical or strategic bias (Lu et al., 2008). Some recent studies estimate passenger preferences from observed route choices using detailed trip-level records. Hörcher et al. (2017) study the Hong Kong Subway and Yap et al. (2020) use Dutch tram and bus data. Both papers estimate crowding costs relative to travel time, so monetary valuation requires an external value of time. Similarly, Hahn et al. (2023) evaluate fare subsidies in BART using crowding valuations drawn from prior research. We instead build an economic model that incorporates price, crowding, and schedule deviation jointly and estimate WTP to avoid crowding in monetary terms using exogenous fare variation generated by the EBD. We further address the endogeneity of crowding with a novel instrument and estimate preference heterogeneity by income through a random-coefficient specification.

Second, we contribute to the literature on optimal pricing under transit congestion externalities. Prior theoretical work characterizes transit pricing and subsidies with crowding (Parry and Small, 2009; de Palma et al., 2017; Coulombel and Monchambert, 2023). To our knowledge, we provide the first empirical estimate of an optimal crowding tax based on monetary crowding valuations identified from exogenous fare variation. We also compare the tax against quantity controls that ration access through waiting rather than price. The welfare comparison connects to the theory of price versus quantity controls (Weitzman, 1974) and empirical work on market and nonmarket allocation of automobile licenses (Li, 2018). Under heterogeneous disutility from crowding, we show that second-degree price discrimination through a two-class configuration elicits self-selection and produces a Pareto improvement, a result that complements Cook and Li (2025) on high-occupancy toll lanes.

Third, we contribute to the broader literature on externalities from shared resource use. The theory of clubs (Buchanan, 1965) and the literature on common-pool resource use (Gordon, 1954; Hardin, 1968) examine how individual use affects others. Empirical studies quantify these externalities in fisheries (Huang and Smith, 2014) and groundwater (Ryan and Sudarshan, 2022). Timmins and Murdock (2007) estimate congestion costs at recreational fishing sites while addressing endogenous site choices. We study in-vehicle crowding, a common experience in daily commuting. Here, passengers compete for space through simultaneous use, and each additional rider reduces the comfort of others. Unlike the depletion of fish or groundwater stocks, this externality takes the form of a loss

in service quality that commuters experience repeatedly. Our estimates quantify this recurring cost and inform the allocation of limited space in a widely used service.

Fourth, we develop new measurement tools that exploit the granular trip-level data and the network structure of the subway system. Although trip-level transit records are increasingly used in empirical work (e.g., Gu et al., 2024), they typically contain only entry and exit times and stations and do not allow researchers to trace passengers through the network. We develop an iterative algorithm that imputes each passenger’s optimal route, real-time location, and in-train crowding. We also propose a novel way of measuring the optimal arrival time, a crucial but typically unobserved input in a large class of structural travel models with endogenous congestion, including the bottleneck model (Vickrey, 1969; Arnott et al., 1993). Our approach uses the arrival distribution of riders unaffected by the EBD to infer EBD riders’ optimal arrival times non-parametrically through quantile matching. We contribute to studies of scheduling flexibility among commuters (e.g., Gu et al., 2024; Kreindler, 2024; Hall, 2024). The same network structure also provides the incidental-companion instrument for crowding. Together, these tools allow us to measure crowding and rescheduling disutility and may prove useful in other settings with similarly granular data.

The rest of the paper is organized as follows. Section 2 describes the institutional setting and the data. Section 3 introduces the demand model and provides an overview of the estimation strategy. Section 4 describes the construction of key variables and instruments and presents summary statistics. Section 5 documents stylized facts that motivate the modeling choices. Section 6 presents the estimation results for both the homogeneous logit and the random-coefficient specifications. Section 7 conducts the welfare and counterfactual analysis. Section 8 concludes.

2 Background and Data

2.1 Background

Beijing is China’s capital and second most populous city. In the first two decades of the 21st century, it has built one of the world’s largest subway systems, almost from scratch. By 2016, the network consisted of 19 lines and more than 340 stations, carrying 3.7 billion annual passenger trips. The subway is especially important for long-distance commuting because of its speed and reliability. Excluding walking trips, the subway accounted for 15% of commuting trips and roughly 40% of commuting passenger mileage in 2014 (Beijing Transport Institute, 2015). A typical subway trip covers about 15 kilometers and takes approximately 37 minutes, including time spent accessing, waiting, and exiting.

The fare for a subway ride is a step function of the track distance between the origin and the destination. A single trip starts at 3 RMB for distances under 6 kilometers, with the fare increasing by 1 RMB for each additional increment of up to 20 kilometers, eventually reaching 10 RMB for

trips between 92 and 112 kilometers. Starting in December 2015, an early-bird discount (EBD) was introduced at 16 stations on two suburban lines, the Batong Line and the Changping Line. Passengers who tap in at one of these stations before 7 AM receive a 30% fare reduction.


The 16 EBD stations are among the busiest during the morning rush, but they are not the only busy stations. Other suburban lines connecting dense residential neighborhoods with the city center exhibit similar ridership patterns. The EBD was later expanded to eight additional stations on Line 6 by the end of 2016, which falls outside our sample period. We select 16 stations on Lines 5 and 6 as comparison stations. Line 6 runs parallel to the Batong Line, and the eight stations we select received the EBD after our sample period. Line 5 cuts through the city center from north to south, with its northern branch paralleling the Changping Line and passing through dense residential neighborhoods. The eight Line 5 stations we choose have been in active discussion for inclusion in the EBD program. Appendix Figure C.2 shows the location of EBD and control stations in the Beijing Subway network.

Figure 1 compares passenger flows at EBD and control stations during morning hours. The two groups exhibit strikingly similar in-train ridership profiles, with both peaking at approximately 15,000 to 16,000 passengers around 8 AM. The similarity supports the use of control stations as a valid comparison group. The figure also reveals the behavioral response to the EBD. At EBD stations, the entry flow displays a visible dent just after 7 AM, with a corresponding bunching of entries just before the cutoff. This pattern is consistent with riders advancing their departures to qualify for the 30% discount. No such discontinuity appears at control stations, where the entry flow rises smoothly through the morning peak. The in-train passenger counts, which cumulate entries from upstream stations, also show dent around the cutoff, albeit less sharp than that for the number of entries. The discontinuities generated by the EBD may seem small and localized in the narrow neighborhood around the time cutoff, but as Figure 2 illustrates, as the EBD affects the crowdedness of subway cars, its impact ripples beyond those who directly take advantage of the discounts.

2.2 Data

The primary data source is trip-level records collected by Beijing Subway’s automated fare collection (AFC) system (Beijing Municipal Institute of City Planning and Design, 2015). The Beijing Subway uses an electronic smartcard as its payment method. The smartcard is popular because it is easy to use and only card users qualify for fare discounts. By comparing official ridership statistics with our data, we estimate that between 90% and 95% of subway trips were paid for by smartcard around the time of the EBD. In our data, each record captures the tap-in and tap-out stations and timestamps for a single trip, and trips can be linked across days through an anonymized card identifier. The main data used in this paper consist of the universe of smartcard trips on 21 working days

Figure 1: Passenger Flows during Morning Hours



Notes: This figure plots the number of passengers (in thousands) at EBD stations (panel a) and control stations (panel b) by 5-min time bin during morning hours. The four series represent passengers entering the station, exiting the station, transferring through the station, and riding in trains passing through the station. The orange dashed vertical line marks the 7 AM EBD cutoff, and the two gray dashed vertical lines mark 6:30 AM and 8:30 AM, which bound the sample period used in the analysis. Both groups of stations exhibit similar in-train ridership levels, peaking around 8:00 to 8:30 AM. At EBD stations, entry passengers display a visible dent just after 7 AM as some riders advance their departures to qualify for the discount, with a corresponding bunching just before 7 AM. No such discontinuity appears at control stations.

in five non-consecutive collection periods in 2016.⁷

We supplement the AFC records with several additional data sources. We geocode all subway stations and collect the track distance and fare for each station pair from Beijing Subway’s website. We collect train frequencies by line and station, as well as the types of passenger cars and the number of cars per train, from which we compute the total floor area of each train. For train floor area, we do not distinguish between seating and standing areas. We also use the 2015 Beijing Household Travel Survey (Beijing Transport Institute, 2015), a large representative survey of approximately 100,000 households that records each respondent’s home and work locations, daily travel details, and transportation mode choices. We further use grid-level commuting flow data from Baidu Maps to construct OD-pair-specific income distributions. The data and our matching procedure are described in Section 4.1.3. Appendix A provides a detailed description of all data sources used in this paper.

3 Model


This section presents a model of travel demand in which subway passengers choose their departure time, trading off the fare, expected in-train crowding, and the cost of deviating from an optimal schedule. We first describe the model and then provide an overview of the estimation strategy.

⁷The collection periods fall in March, June, September, November, and December 2016. June 9–10 and September 15–16 are excluded because of the Dragon Boat Festival and Mid-Autumn Festival holidays, respectively.

3.1 Model of Subway Demand

Figure 2 illustrates the intuition behind our estimation. For simplicity, suppose the density of travel demand is flat over time, as indicated by the horizontal gray dashed line. Now consider the introduction of the EBD before the cutoff time t_0 . Take a passenger who originally planned to depart right after 7 AM, at t^0 . Once the discount is in place, they have an incentive to tap in before the cutoff to claim the lower fare. If the cost of rescheduling increases with the distance from the original departure time, they have no reason to move any earlier than necessary and would depart just before the cutoff, at t^1 . However, other passengers may reason the same way, making the train right before the cutoff extremely crowded. Given how much they are willing to pay to avoid crowding, the passenger may instead choose to depart at an earlier time, at t^2 . The EBD can also generate ripple effects on those who do not directly benefit from the discount. Consider another passenger who planned to leave at $t^{0'}$. They may now choose to depart earlier, at $t^{1'}$, not to qualify for the EBD but to take advantage of the less crowded subway cars. The red curves trace the likely profile of trip density in the presence of the EBD. The network structure of the subway further complicates passengers' calculations, as those who respond to the EBD generate spillover effects that propagate through the network.

Figure 2: The Departure Time Trade-off under the Early-Bird Discount



Notes: This figure illustrates passengers' departure time choices under the early bird discount.

We model passengers' departure-time choices, taking their origin–destination (OD) pair and route as given. Departure times are discretized into 15-minute bins between 6:30 and 8:30 AM, measured by the tap-in time at the origin station. We define a *market* as an OD pair j on a calendar date d , and each departure-time bin $t \in \mathcal{T}$ is a *product*. The outside option represents traveling between the same origin and destination during morning rush hours by other modes.

Passenger i 's utility from departing at time t in market (j, d) is

$$U_{ijdt} = V_{ijdt} + \varepsilon_{ijdt}, \quad (1)$$

where the error term, ε_{ijdt} , is i.i.d. following a Type-I extreme value distribution, and the deterministic

component is

$$V_{ijdt} = \alpha_i P_{jdt} + \beta_i \mathbb{E}_d[\text{Crowd}_{jdt}] + \rho_i \left| t + \mathbb{E}_d[\text{TT}_{jdt}] - t_{ijd}^{OA} \right| + \kappa_{jd} + \phi_{dt} + \xi_{jdt}. \quad (2)$$

The fare P_{jdt} varies across OD pairs as a step function of distance and varies across time bins within an OD pair at EBD stations, where passengers receive a 30% discount for tapping in before 7 AM.

On each date d , passengers form expectations about crowding and travel time before choosing when to depart. The operator $\mathbb{E}_d[\cdot]$ denotes these expectations. The term $\mathbb{E}_d[\text{Crowd}_{jdt}]$ is the expected total in-train crowding experienced during the trip, measured in passengers $\cdot$ minutes per square meter; it aggregates crowding density over the duration of the passenger's in-train exposure. We proxy expected crowding with $\widehat{\text{Crowd}}_{j,\text{dow},t}$, where dow denotes the day of the week corresponding to date d . These fitted values come from a regression of realized crowding on OD-by-day-of-week-by-time-bin fixed effects, equivalent to averages across dates within each cell. They capture the systematic conditions that experienced commuters can anticipate.

The rescheduling term $\left| t + \mathbb{E}_d[\text{TT}_{jdt}] - t_{ijd}^{OA} \right|$ captures the absolute deviation between the passenger's expected arrival time—the sum of departure time t and expected travel time—and their optimal arrival time t_{ijd}^{OA} . We proxy expected travel time with $\widehat{\text{TT}}_{j,\text{dow},t}$, using the same averaging procedure as for crowding. We assume the penalty from deviating from the optimal arrival time is symmetric and linear in time.⁸

The remaining terms are OD-by-date (market) fixed effects κ_{jd} , which absorb the common value of completing trip j on date d ; date-by-time-bin fixed effects ϕ_{dt} , which capture departure-time patterns common across OD pairs on each date; and a structural demand shock ξ_{jdt} that varies at the OD-by-date-by-time-bin level and is unobserved by the econometrician. The coefficients α_i , β_i , and ρ_i represent the marginal (dis)utility of fare, crowding, and schedule deviation, respectively.

First consider the simple case where preference is homogeneous: $\alpha_i = \alpha$, $\beta_i = \beta$, and $\rho_i = \rho$. Normalizing the outside option utility to zero, the market share of time bin t in market (j, d) takes the standard logit form:

$$S_{jdt} = \frac{\exp(\delta_{jdt})}{1 + \sum_{\tau \in \mathcal{T}} \exp(\delta_{jd\tau})}, \quad (3)$$

where $\delta_{jdt} \equiv \alpha P_{jdt} + \beta \widehat{\text{Crowd}}_{j,\text{dow},t} + \rho \left| t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{OA} \right| + \kappa_{jd} + \phi_{dt} + \xi_{jdt}$ is the mean utility, where t_{jdt}^{OA} is the mean schedule deviation of passengers choosing the same product. Following Berry (1994), the share equation can be inverted to yield the linear estimating equation:

$$\log S_{jdt} - \log S_{jd0} = \alpha P_{jdt} + \beta \widehat{\text{Crowd}}_{j,\text{dow},t} + \rho \left| t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{OA} \right| + \kappa_{jd} + \phi_{dt} + \xi_{jdt}, \quad (4)$$

⁸The symmetry and linearity assumptions can be relaxed. We maintain them here for two main reasons. First, the EBD incentivizes riders to reschedule to an earlier departure time. Second, as we show below, the deviation in travel time is small in general, so allowing for nonlinear effects does not make a substantial difference.

Because $\log S_{jdt}$ varies only at the market level, it is absorbed by κ_{jd} . Price, the expected crowding experience, and the expected schedule deviation all enter the utility function in Equation (1) in levels, so ratios β/α and ρ/α give the marginal willingness to pay (MWTP), in Chinese yuan (RMB), to avoid one additional passenger per square meter for one minute and the MWTP to reduce schedule deviation by one minute, respectively.

Passengers with different income levels may have different price elasticities and derive different levels of disutility from crowding and rescheduling. To capture this heterogeneity, we allow preferences to vary with income:

$$\alpha_i = \alpha_0 + \alpha_1 I_i, \quad \beta_i = \beta_0 + \beta_1 I_i, \quad \rho_i = \rho_0 + \rho_1 I_i, \quad (5)$$

where I_i is the passenger income, recentered at the sample mean. Because the smartcard records do not contain passenger income, we simulate income from an OD-specific distribution of potential commuters.⁹ With heterogeneous preferences, the utility decomposes into a mean component δ_{jdt} , common to all passengers in the market, and an individual deviation μ_{ijdt} that depends on income:

$$\begin{aligned} V_{ijdt} = & \underbrace{\alpha_0 P_{jdt} + \beta_0 \widehat{\text{Crowd}}_{j,\text{dow},t} + \rho_0 \left| t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{OA} \right| + \kappa_{jd} + \phi_{dt} + \xi_{jdt}}_{\delta_{jdt}} \\ & + \underbrace{\alpha_1 I_i P_{jdt} + \beta_1 I_i \widehat{\text{Crowd}}_{j,\text{dow},t} + \rho_1 I_i \left| t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{OA} \right|}_{\mu_{ijdt}}. \end{aligned} \quad (6)$$

The random-coefficient model no longer allows for a closed-form expression for market shares, and we estimate it using the conventional iterative procedure in the BLP literature (Berry et al., 1995): for a given vector of non-linear parameters $\theta = (\alpha_1, \beta_1, \rho_1)$, predicted shares are computed by simulation over income draws, the mean-utility vector δ is recovered via the contraction mapping, and the resulting structural residuals are interacted with a set of instruments to form the GMM objective.

3.2 Estimation Strategy Overview

Estimating the model requires measuring crowding and optimal arrival times and identifying passengers' responses to fares, crowding, and rescheduling. We first describe how we construct the unobserved variables, then explain the variation and assumptions used to identify the preference parameters. Section 4 presents the detailed measurement procedures and instrument construction.

⁹We use Baidu Maps commuting flows and income brackets to construct an income distribution for potential commuters in each OD pair. We fit a beta distribution and draw 150 incomes for each OD-date market. These draws represent potential commuters rather than only observed subway riders; the model allows their choice between subway travel and the outside option to depend on income. Section 4.1.3 describes the construction.

3.2.1 Measurement

The automated fare collection (AFC) records cover subway trips throughout the network, recording each passenger’s origin and destination stations and tap-in and tap-out times. They do not reveal which route a passenger takes, which train they board, or the crowding they experience. We therefore reconstruct passengers’ routes and locations during their trips. We develop an iterative algorithm that jointly estimates the optimal route from the origin to the destination and the time spent walking, waiting, and transferring at each station, with in-train time imputed from the published timetable. Dividing the estimated number of passengers sharing a train by its available floor area yields crowding density. Summing this density across trip segments, weighted by time spent on each segment, gives total in-train crowding exposure.

The optimal arrival time is also unobserved. Passengers choose departure times partly based on their preferred arrival times, and omitting this variable would confound the crowding coefficient: OD-time combinations with clustered optimal arrival times would exhibit both high crowding and high unobserved demand, biasing the estimate of β . We infer each passenger’s optimal arrival time from “coworkers”—subway passengers who arrive at the same destination around the same time on the same date but originate from stations not directly affected by the EBD. Their arrival distribution reflects arrival choices without the direct fare incentive to shift schedules. Within each destination-date cell, we reweight this distribution to match the income composition of EBD riders and assign each rider the coworker arrival time at the same percentile of the reweighted distribution. Section 4.1.2 describes the procedure.

Both measurement strategies exploit the *network structure* of the transit system. The train-assignment algorithm traces passengers across the subway network to compute crowding on shared segments. The coworker-based imputation identifies riders who converge at the same destination from different parts of the network. As we discuss next, the same network structure also underpins our instrumental variable strategy for crowding.

3.2.2 Identification

In our setting, identification of the fare coefficient α relies on variation induced by the early-bird discount. The policy reduces fares by 30% for trips originating from 16 designated stations before 7 AM. Outside the discount, fares P_{jdt} depend on the track distance between the origin and destination stations. The EBD stations were not randomly selected; they are among the busiest suburban stations. We include OD-by-date fixed effects κ_{jd} to absorb differences in demand levels across OD pairs on each date, so α is identified from within-market variation over departure-time bins. We also include comparison stations on two other suburban lines that exhibit similar ridership patterns but were not covered by the EBD. The comparison stations help identify the date-by-time-bin fixed effects ϕ_{dt} , allowing us to distinguish the response to the EBD fare reduction from changes in demand over

the morning.¹⁰ Identification requires that, conditional on the fixed effects and other trip attributes, unobserved demand ξ_{jdt} does not vary systematically with the EBD-induced fare reduction.

Expected crowding $\widehat{\text{Crowd}}_{j,\text{dow},t}$ may remain positively correlated with ξ_{jdt} , even after controlling for fares, rescheduling costs, and the fixed effects. This demand shock varies jointly across OD pairs j , dates d , and departure-time bins t . A subway ride in a particular OD pair at a particular time is crowded not because passengers enjoy crowded trains, but because traveling in that OD pair at that time has high unobserved value *despite* crowding. For instance, if many residents near the origin station also work near the destination and conventional work hours require arrival by 9 AM, commuting demand concentrates at specific departure times. To the extent that the fixed effects and measured rescheduling costs do not capture these OD-specific departure-time preferences, they generate both high crowding and high unobserved demand. This positive correlation biases the OLS estimate of β toward zero, or can even make it positive, understating the true disutility of crowding.

We instrument for crowding using passengers whose trips respond to different demand shocks but overlap with the focal trip on shared train segments. These passengers travel from other origins to distant destinations, adding to crowding on the shared segments. We call them *incidental companion passengers* (ICPs). We restrict ICPs to infrequent riders and assume that, conditional on the model’s controls, their demand shocks are unrelated to the focal trip’s demand shock ξ_{jdt} . Our instrument is expected incidental companion time (ICT), which weights each ICP by the duration of shared travel and uses the same expectation procedure as crowding. Section 4.2 describes its construction.

To identify the rescheduling coefficient ρ , we assume that, conditional on the model’s controls, $|t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{\text{OA}}|$ is uncorrelated with ξ_{jdt} . One concern is that unobserved demand may increase passenger volume and travel time, changing the rescheduling term. We show below that in-train crowding does not meaningfully slow trains and that controlling for in-station crowding does not materially change the crowding valuation estimates. Another concern is that EBD and control riders may differ in their preferred arrival times t_{ijd}^{OA} even for the same destination and date. These differences can affect both the imputed schedule deviation and demand for particular OD-time combinations, inducing correlation with ξ_{jdt} beyond what the fixed effects absorb. We show in Section 4.1.2 that pre-policy data reveal no systematic differences in arrival patterns between EBD and control riders associated with the 7 AM cutoff.

¹⁰The demand model includes price level, so using trips only from EBD stations, the model can still be estimated with OD-by-date and date-by-time-bin fixed effects. But the model is only identified due to the specific functional form. For example, with only EBD stations, the model is unidentified with log price.

4 Measurement and Identification

4.1 Measurement

4.1.1 Optimal Route, Train Assignment, and Crowding

To measure in-train crowding, we need to determine each passenger’s location at every point during their trip. There are two key challenges to this task. First, the AFC data record only each passenger’s tap-in and tap-out stations and times, not which trains they board. We also do not observe trains’ real-time locations or the crowding density on platforms or inside cars.¹¹ Second, with 19 lines and more than 340 stations, the Beijing Subway is a complex network, and many OD pairs admit multiple feasible routes. We need to determine which route each passenger takes.

Optimal route and train assignment. We assume that passengers take the route with the shortest travel time and identify optimal routes using observed data. Specifically, we define a route $\mathcal{P}_{jdt}$ as an ordered sequence of stations and their associated times,

$$\mathcal{P}_{jdt} = \{S^O(t^O), S_1(t^1), \dots, S_n^R(t^n), \dots, S^D(t^D)\},$$

where S^O and S^D denote the entry and exit stations, respectively. S_n denotes pass-through stations in the route. S^R denotes a transfer station if applicable. We let $\mathcal{R}_{\mathcal{P}_{jdt}}$ represent the set of transfer stations along route $\mathcal{P}_{jdt}$. t in the parentheses indicates the time when the passenger reaches the corresponding station. We observe entry time (t^O) and exit time (t^D), but not the en-route stations or time arriving at and departing from those stations (t^1, t^2 , and so on).

The total travel time along route $\mathcal{P}$ can be decomposed as

$$TT_{jdt}^{\mathcal{P}} = \underbrace{ST_{jdt}^{\mathcal{P}}}_{\text{in-train}} + \underbrace{AT_{t^O}^O}_{\text{entry}} + \underbrace{\sum_{R \in \mathcal{R}_{\mathcal{P}}} RT_{t^R}^R}_{\text{transfer}} + \underbrace{ET_{t^D}^D}_{\text{exit}}. \quad (7)$$

The total travel time can be calculated from the tap-in and tap-out timestamps. $ST_{jdt}^{\mathcal{P}}$ is the in-train time, defined as the sum of scheduled train travel times between adjacent stations along the route. A key assumption we maintain throughout the paper is that trains always run on time. We validate this assumption, especially whether crowding slows down the train, later in the paper. Therefore, once the route is known, $ST_{jdt}^{\mathcal{P}}$ can be calculated from published train timetables. $AT_{t^O}^O$ is the entry time between tapping into the origin station and boarding the first train; $ET_{t^D}^D$ is the time between leaving the train at the destination station and passing through the exit fare gate; and $RT_{t^R}^R$ is the

¹¹Most trains in the Beijing Subway are equipped with real-time location devices, which are used to display estimated arrival times at stations. Platforms and passenger cars also have cameras that could potentially be used to derive crowding measures. However, we do not have access to those data.

transfer time at transfer station R when arriving at time t^R . We observe when passengers enter and leave the fare gates, but not when they board, transfer between, or leave trains. We therefore infer the entry, transfer, and exit components of total travel time.

We jointly estimate the optimal route and entry, exit, and transfer times through an iterative procedure applied to the entire subway network. We begin by assigning initial guesses for each in-station component. These initial values allow us to back out arrival times at transfer stations and calculate the total travel time for each candidate route. Using the calculated travel times, we apply Dijkstra’s algorithm to identify the optimal route for each OD pair.

We update the estimates of in-station time components by regressing the observed median total travel time $\hat{T}T_{jdt^O}$ for trips in each OD-tap-in-time bundle on in-train time $\hat{S}T_{jdt^O}$ and a set of station-by-purpose (entry, exit, transfer) specific fixed effects capturing in-station time cost:

$$\hat{T}T_{jdt^O} = \hat{S}T_{jdt^O} + AT_{t^O}^O + \sum_{R \in \mathcal{R}} RT_{t^R}^R + ET_{t^D}^D + \varepsilon_{jdt^O}. \quad (8)$$

Note that a station can serve as an entry station, an exit station, or a transfer station for different trips. We allow for a different time cost in that station-time-bin bundle for each purpose. The regression is estimated using the ordinary least squares (OLS), pooling across the entire subway network and weighting by the number of trips in each jdt^O . The number of unique OD pairs is orders of magnitudes larger than the number of unique station-purpose pairs, so there are enough degrees of freedom to flexibly identify entry, exit, and transfer times specific to each station. We recalculate the optimal route and in-train time using the updated estimates, and iterate until all in-station time parameters and the implied route choices converge.¹²

Given a trip’s optimal route, Equation (8) faces a collinearity problem. The total travel time is observed, the in-train time is determined by the timetable, and the error term is mean-zero, so the in-station time components must sum to a constant equal to the mean total travel time minus in-train time. To break the collinearity, we normalize the mean exit time to two minutes.

Appendix Table C.1 summarizes the estimated in-station times. Entry time, from tapping in at the origin station to boarding the train, averages 3.5 minutes with a standard deviation of 1.84 minutes. This variation may reflect station configurations, train frequencies, and in-station crowding. Exit time averages two minutes by construction, and has a standard deviation of 0.9 minutes. Transfer time, at an average of 4.43 minutes, is the longest in-station component, with some transfer stations requiring more than six minutes. The last column of the table reports the number of observations used to estimate each component. Because all trips have an entry and exit station, these two components are estimated from approximately 33.4 million trips. The average trip involves slightly more

¹²While the optimal route may in principle vary over the course of the day due to time-varying in-station congestion, in practice the fastest route for a given OD pair rarely changes by time of day. This is because the second-best routes typically cost substantially more time. We therefore assign the most common optimal route across all departure times to each OD pair.

than one transfer, so transfer time is estimated from 34.8 million trips.

Not all passengers traveling from O to D with tap-in time t^O arrive at the same time. There is a distribution of total travel time, which we attribute to idiosyncratic factors in the in-station segments. For example, some passengers walk faster than others, some are lucky to catch a departing train while others must wait for the next one, and some may stop at a convenience store or vending machine inside the fare gates. While assuming all passengers take the same optimal route, we treat the sources of travel time variation as agnostic and proportionately scale each in-station segment of the median passenger's time cost to fit the observed total time for each individual trip. For example, suppose the imputed travel time for a trip in jdt^O is 40 minutes, composed of 5 minutes entry, 10 minutes transfer, 20 minutes in-train, and 5 minutes exit. If an actual trip in that bundle takes 44 minutes (10% longer than imputed), we assign 6 minutes for entry, 12 minutes for transfer, and 6 minutes for exit, while the in-train time remains fixed at 20 minutes.

The route assignment algorithm rests on several assumptions that merit discussion. First, train speed is assumed not to vary with crowding. During the study period, the Beijing Subway was largely new and equipped with advanced computerized control systems. Stopping times at stations are automated, and most stations have platform screen doors to prevent accidents. Unlike some older systems, delays were rare. We show in Section 5 that the empirical relationship between crowding and travel time is negligible. Second, we assume that passengers do not detour to avoid crowded routes. Detours typically cost substantially longer time, and Section 5 provides evidence that route-switching to avoid crowding is uncommon. Third, we assume that the distribution of travel time within an OD-time bundle reflects idiosyncratic shocks and is unrelated with crowding levels.

Having decomposed each passenger's route and time, we can determine their real-time location at any point during the trip. We define a *segment* s as the track between two adjacent stations. Each passenger's trajectory is discretized into segment-by-five-minute cells. When a passenger traverses multiple segments within a single five-minute bin, their presence is distributed across segments in proportion to the time spent on each. When a single segment spans multiple five-minute bins, it is partitioned into sub-segments assigned to the corresponding bins. Summing over all passengers yields $N_{s\tau}$, which is the weighted ridership on segment s during time bin τ .

Crowding density and experience. We focus on in-train crowding.¹³ We define the in-train *crowding density* $k_{s\tau}$ as the number of passengers per square meter on segment s in time bin τ , obtained by dividing $N_{s\tau}$ by the total floor area of a six-car train.¹⁴ $k_{s\tau}$ reflects the degree of crowding


¹³Passengers may also derive negative utility from crowded stations. Our route assignment algorithm also allows us to calculate the number of passengers in each station at any given time. We focus on in-train crowding because it is arguably the most important dimension of the disamenity. Including in-station crowding is possible, but because the station is much larger and passenger distribution is uneven, relevant crowding density is difficult to impute. We show in robustness checks that including in-station crowding does not affect the estimates of the WTP to avoid in-train crowding.

¹⁴Each train in our sample consists of six passenger cars measuring 19 meters long and 2.8 meters wide, providing a total floor area of approximately 320 square meters. We do not subtract floor space occupied by seats and other on-train

a passenger experiences at a given time and has units of persons per square meter.

For each passenger i , the in-train *crowding experience* K_i aggregates the crowding density over all segment-time bins along their route, weighted by the time spent in each. The in-train crowding experience has units of person-minutes per square meter, capturing the total discomfort as the product of the intensity and the duration of crowding exposure. These trip-level exposures underlie the expected crowding proxy used in the demand model. We aggregate K_i to the product level by computing the mean across all trips within each OD-by-date-by-15-minute-bin (jdt) cell. Figure 3 illustrates this measure: the area under the density profile is the passenger’s cumulative in-train crowding exposure.

Figure 3: Crowding Density and Cumulative Exposure



Notes: Illustrative example, not sample estimates. The passenger spends 10 minutes at a density of 2 persons per square meter and 20 minutes at a density of 5. The shaded areas sum to 120 person-minutes per square meter of crowding exposure. Over the 30-minute trip, average density is 4 persons per square meter.

4.1.2 Optimal Arrival Time and Rescheduling


The rescheduling cost in the demand model requires measuring the optimal arrival time t_{ijd}^{OA} for each passenger. This variable is fundamentally unobserved. Observed arrival times may not be optimal because passengers trade off schedule deviation against fare savings and expected crowding. In particular, the EBD incentivizes passengers to depart earlier than they otherwise would, so observed

equipment, nor do we model the probability of finding a seat or differential utilities from sitting and standing. Using smartcard data from the Hong Kong subway similar to ours, Hörcher et al. (2017) presents an algorithm to determine the probability of finding a seat. We do not observe detailed train headways by line and time of day. The ratio of segment run time to headway gives the average number of trains on a segment. Approximate morning-peak headways suggest a ratio of about 0.8–1.2, so using one train’s floor area approximates average on-train density when this ratio is close to one.

arrival times of trips from EBD stations are also likely to be earlier than optimal. Exploiting the network structure of the subway system, we construct the optimal arrival time using trips that arrive at the same destination station but originated from non-EBD stations.

The key idea is that passengers arriving at the same destination in the morning peak hours as the focal passenger are likely responding to the same need for travel. We refer to these passengers as “coworkers” because they travel to the same destination on the same date and may face similar workplace schedules. Their arrival distribution reflects arrival choices without the direct fare incentive to shift departure times. If control station riders travel to the same destination on the same day as EBD riders, the two groups face similar workplace constraints and destination-side demand conditions. The control riders’ arrival distribution therefore serves as a natural counterfactual for what EBD riders would have chosen absent the discount.

Figure 4: Optimal Arrival Time Imputation and Validation



Notes: Panel (a) illustrates the quantile matching procedure. The blue distribution shows observed arrivals at destination D from EBD stations, shifted earlier by the fare discount. The gray distribution shows arrivals from control stations traveling to the same destination on the same day. Each EBD rider is matched to the control distribution at the same within-group percentile rank, yielding the imputed optimal arrival time t_i^{OA} . Panel (b) plots the mean signed deviation (actual minus optimal arrival time, in minutes) for EBD station trips by 15-minute entry time bin. Negative values indicate that passengers arrived earlier than the imputed optimal time. The dashed line marks the EBD discount cutoff at 7 AM.

Figure 4 panel (a) provides a graphical illustration of this approach. For a given date d and destination station D , we collect all trips arriving at D during the morning rush from control stations. This set of trips forms the control arrival distribution. We then collect the corresponding trips from EBD stations. We further refine the control arrival distribution by reweighting it to match the income composition of EBD riders in the same date-destination cell. Because the EBD discount shifts some passengers to earlier departures, the EBD arrival distribution is likely shifted to the left relative to the control distribution, as illustrated in Figure 4 panel (a). We perform quantile matching to map

each EBD rider to the control distribution. Specifically, let p_i denote the percentile rank of EBD rider i within the EBD arrival distribution for that date-destination cell. The imputed optimal arrival time is

$$t_i^{OA} = F_{\text{Ctrl}}^{-1}(p_i), \quad (9)$$

where F_{Ctrl}^{-1} is the inverse cumulative distribution function of the income-reweighted control arrival distribution. The matching is performed at high precision using continuous seconds. Cells with fewer than ten EBD trips or fewer than ten control trips are excluded from the procedure.

For trips originating from control stations, we set the optimal arrival time equal to the actual arrival time, $t_{\text{Ctrl}}^{OA} = t_{\text{Ctrl}}^{\text{Actual}}$, so that rescheduling cost is zero by construction. This ensures that any nonzero rescheduling cost in the estimation comes from the EBD group, where the fare discount induces measurable departures from schedules not directly unaffected by the EBD.

This approach has several advantages. It relies entirely on contemporaneous data, eliminating concerns about temporal shifts in commuting patterns that would arise from using pre-policy historical records. By matching within the same date-destination cell, it captures date-specific demand conditions such as weather or local events that affect both EBD and control riders symmetrically. This approach is essentially non-parametric. It assigns coworkers through matching two *distributions*, rather than directly matching individuals. Without knowing the exact workplace for each rider, assigning coworkers to each rider would require more assumptions and likely introduces additional biases.

A key assumption underlying the method is that passengers from EBD and control stations traveling to the same destination face similar optimal arrival incentives. Appendix Figure C.3 provides supporting evidence by comparing the destination distributions of the two groups. Of the 241 unique destination stations in the sample, 188 are served by riders from both EBD and control origins. The remaining 53 single-group destinations account for a negligible share of total ridership (0.55% for EBD-only destinations and 1.21% for control-only destinations). The Pearson correlation between the two destination share vectors is 0.46, and the Spearman rank correlation is 0.68, indicating that the two groups share broadly similar destination patterns despite originating from geographically distinct lines.

Figure 4 panel (b) validates the imputed rescheduling measure by plotting the mean signed deviation (actual minus optimal arrival time) for trips originated from EBD stations across 15-minute entry time bins. Negative signed deviations indicate that the passenger arrived earlier than optimal, which is the expected direction for EBD riders who advance their departures to qualify for the discount. Riders who tap in before the 7 AM cutoff arrive approximately one to two minutes earlier than their imputed optimal time. After 7 AM, when the discount no longer applies, deviations become smaller. The largest deviations are observed for trips that depart around the cutoff. The clear discontinuity at the cutoff confirms that the imputed optimal arrival time captures meaningful variation

in rescheduling behavior. The vast majority of the deviations calculated for EBD trips are negative. The rescheduling variable used in the demand model is the absolute deviation between the expected arrival time and the imputed optimal arrival time, $|t + \widehat{\text{TT}}_{j,\text{dow},t} - t_{jdt}^{OA}|$.

Pre-existing differences between EBD and control riders' optimal arrival-time distributions could cause the imputed deviation to reflect scheduling preferences formed before the policy as well as a response to the discount. Applying quantile matching to 2015 data, before the EBD, yields a smooth deviation profile at future-EBD stations with no discontinuity at 7 AM. Comparing the 2016 and 2015 profiles shows an additional advance of 3.3 and 2.9 minutes in the two pre-7 AM bins. This comparison supports interpreting the change near the cutoff as policy-induced rescheduling.

4.1.3 Income Distribution

Passengers may have heterogeneous preferences for trip characteristics. We are particularly interested in preference heterogeneity by income levels, which has implications for the distributional impacts of crowding and crowding-reducing policies. The smartcard data do not have information on user demographics. We match passenger income using novel outside data sources.

We use grid-level commuting flow data provided by Baidu Maps, a leading provider of digital map services in China. Baidu Maps provides location service on smart devices and is embedded in many mobile applications. It is able to track millions of smart mobile devices (mostly smartphones) with precise and high-frequency location data. In particular, it determines a device's usual night-time location, which it refers to as home, and its usual daytime location, referred to as work. The commuting flow data contains the number of commuters between all pairs of home and work locations in Beijing, where each location is a small grid with an edge of approximately 700 meters, and thus an area of about 0.5 square kilometers. The data are based on the usual location patterns in the three months ending in November 2021 and cover 7.6 million unique devices with home locations distributed across roughly 15,000 grids.¹⁵ Importantly, the data reports the share of commuters in each home-work location pair in five monthly income brackets. User income is inferred by Baidu Maps based on the information it has access to on the device using a machine learning algorithm. The information includes the make and model of the device, applications installed on the device, some information users voluntarily disclosed to Baidu and its partnering applications, and detailed location information. We do not have access to those underlying user characteristics, the exact algorithm Baidu uses, or the training data it uses to validate the imputations. Because there is no official statistics available in China that provide income information at granular geographic levels, the Baidu data is arguably the best such data available. We use the imputed income distribution to estimate the heterogeneous willingness-to-pay to avoid crowding and rescheduling. However, results based on such imputed data should be interpreted with caution.

¹⁵Beijing covers a land area of 16,000 square kilometers and had a population of about 22 million in 2016. The grids do not cover the mountainous regions where few people live and bodies of water.

We assign commuters from the Baidu data to subway station pairs. We define a station’s catchment as an area within a three-kilometer radius. Data from the 2015 Beijing Household Travel Survey (BHTS) (Beijing Transport Institute, 2015) show that very few people choose to ride the subway if they live or work more than three kilometers away in straight-line distance from a station. Catchment areas of different stations can overlap. Users in multiple catchment areas are probabilistically assigned to stations with weights equal to the inverse of distance. For example, suppose a worker lives in the catchment areas of both stations S_1 and S_2 , the distance from their home to S_1 is 1 km while that to S_2 is 0.5 km, then we assign the worker to S_1 with a probability of $1/3$ and to S_2 with a probability of $2/3$. Similarly, the worker’s workplace can be probabilistically assigned to a station. Applying the same algorithm to all commuters in the Baidu data, we define the user pool of each station pair OD . The number of potential commuters is not equal to the number of trips, as people do not necessarily commute every day and there are non-commuting trips. We use the number of subway trips observed in the smartcard data and the share of subway trips from the BHTS data to convert the user pool into market sizes.

For each OD pair, we estimate a continuous income distribution from the grouped income data by fitting a beta distribution via grouped maximum likelihood estimations. The income data is from 2021, while we deflate the incomes to the 2014 price levels, we are still essentially assuming no real income growth and the OD-specific income distribution is time-invariant. For the random-coefficient demand model we estimate below, we draw 150 simulated passenger incomes from the fitted distribution for each OD-date market.

Appendix Figure C.4 plots the distribution of simulated monthly incomes. The weighted mean income is approximately 8,770 RMB per month. The median monthly income of potential subway riders is about 8,000 RMB, and the 90th percentile is around 18,000 RMB. Those numbers are broadly consistent with official statistics. The municipal statistics bureau reported an average monthly salary of 7,706 RMB in 2016. Appendix Figure C.5 maps the mean monthly income by residential and workplace location. High-income residents are spread across suburban areas, while high-income workplaces concentrate along the central business corridors. These geographic patterns are consistent with conventional knowledge of spatial residential and job distributions in Beijing.

We estimate a multinomial logit mode-choice model on the BHTS with individual income and trip distance as predictors, and apply the fitted shares to the Baidu commuting population. Appendix Figure B.1 reports mode choices as functions of income and travel distance. Subway shares rise sharply with trip distance, especially among those with high income. More than 40% of commuters earning above the sample median (8,000 RMB) choosing the subway for commutes longer than 12 km. Cars are primarily used by the high-income for trips of intermediate distances, while bus is used mostly for short trips. Appendix B provides the full estimation results and the predicted mode shares by income and distance.


4.2 Instrument for Crowding

Crowding may remain positively correlated with unobserved demand after accounting for the model’s controls. OD-by-date fixed effects κ_{jd} absorb demand conditions specific to each OD pair on each date, while date-by-time-bin fixed effects ϕ_{dt} absorb departure-time patterns common across OD pairs, including the morning rush. The remaining concern is the residual demand shock ξ_{jdt} , which varies jointly across OD pairs, dates, and departure-time bins. For example, an unusually early work schedule at a destination can increase both demand for a particular OD-time combination and crowding on the trains serving it. This variation reflects demand for the trip as well as passengers’ responses to crowding.

We therefore look for crowding generated by passengers whose origins and destinations expose them to different demand conditions. Consider an office worker traveling to work and an occasional rider traveling from an upstream station to a distant tourist attraction. Their trips share part of the subway route, but their reasons for traveling differ. Variation in visitors’ travel demand can therefore change the crowding experienced by the office worker without changing the worker’s underlying demand for that trip.

We use the subway network to identify passengers whose trips overlap with the focal trip but originate elsewhere and end at distant destinations. We call these riders *incidental companion passengers* (ICPs). Figure 5 illustrates this overlap. Passenger i travels from O to D on date d at tap-in time t . Passenger i' boards upstream at O' and travels to a distant destination D' . The two trips overlap between O and the transfer station R , where passenger i' takes a different line. The shared travel time, $m^{i,i'}$ minutes, is passenger i' ’s contribution to incidental companion time (ICT).

Figure 5: Incidental Companion Passengers



Notes: This figure illustrates the notion of incidental companion passengers (ICPs) and incidental companion time (ICT). Passenger i travels from O to D . Passenger i' boards upstream of O and shares part of i ’s train segment. Passenger i' is an ICP of Passenger i . The time i and i' spend together, between O and the transfer station R , is i' ’s contribution to ICT.

To reduce the likelihood that ICPs share the focal passengers’ regular commuting patterns, we restrict them to infrequent subway users. We define these as passengers with no more than three recorded subway trips during the corresponding collection week. We count trips throughout the day, including recorded weekend and holiday trips, before applying the workday and morning-peak restrictions used in the demand estimation. Our preferred definition also requires ICPs to board

upstream of O and travel to a destination at least eight kilometers from D . Infrequent riders account for only 6.7% of trips in the estimation sample, which consists predominantly of frequent subway users (Table 1).

A focal passenger can have many ICPs. We define ICT as the weighted sum of qualifying ICPs, where the weight is the travel time shared with the focal passenger. Let $\mathcal{I}_i'$ be the set of passenger i 's ICPs. Then

$$ICT_{ijdt} = \sum_{i' \in \mathcal{I}_i'} m^{i,i'}, \quad (10)$$

where $m^{i,i'}$ is measured in minutes. An ICP who shares a longer part of the trip contributes more to ICT than one who overlaps only briefly.

Our instrument is expected ICT, $\widehat{ICT}_{j,\text{dow},t}$, constructed using the same averaging procedure as expected crowding and travel time in Section 3.1. We first calculate ICT for each OD-date-departure-time cell, then average across dates with the same day of the week, holding the OD pair and departure-time bin fixed.¹⁶

The identifying assumption is that, conditional on the model's controls, expected companion demand $\widehat{ICT}_{j,\text{dow},t}$ is unrelated to the focal market's unobserved demand ξ_{jdt} . The spatial and frequency restrictions are intended to limit this correlation. Distant destinations reduce the likelihood of shared workplace schedules or local demand shocks, while infrequent riders are less likely to share the focal passengers' regular commuting patterns. Boarding upstream also makes the companion's trip less directly tied to conditions at the focal origin.¹⁷

Section 6.2 shows that the crowding valuation estimates stabilize as we exclude ICPs whose destinations are close to the focal destination. Imposing an additional distance restriction on ICP origins produces similar estimates. These patterns support the spatial restrictions used to construct the instrument. The first-stage results establish that ICT remains a strong predictor of crowding after these restrictions.

The ICT instrument is structurally analogous to the IV strategy in Hortaçsu et al. (2024), who instrument for endogenous airline fares using the opportunity cost of seat capacity driven by demand from other itineraries sharing the same flight leg. In both settings, the identifying variation comes from travelers who share the same physical infrastructure but serve distinct markets. The exclusion restriction requires that demand from these companion travelers affects the focal outcome only through the endogenous variable.

In addition to expected ICT, we use four instruments based on the EBD policy. The first interacts an indicator for an EBD origin with an indicator for departure before 7 AM. The remaining three

¹⁶In implementation, ICT is divided by train floor area before averaging, so the instrument is measured in person-minutes per square meter. This constant scaling does not change the identifying variation.

¹⁷These restrictions do not require infrequent riders to know nothing about crowding, and do not by themselves establish the exclusion restriction.

interact the EBD-origin indicator with linear, quadratic, and cubic functions of departure time:

$$\text{Treat}_j \times 1\{t < 7:00\}, \quad \text{Treat}_j \times t, \quad \text{Treat}_j \times t^2, \quad \text{Treat}_j \times t^3.$$

Here Treat_j indicates an EBD origin and t denotes departure time, as in the demand model. The first instrument captures the discount cutoff, while the remaining instruments capture differences in the departure-time profile between EBD and control origins.

We control for fare P_{jdt} , so these instruments exploit variation in expected crowding that remains after accounting for fare and the model’s other controls. Figure 2 illustrates why such variation arises. The discount draws passengers toward departures just before 7 AM, increasing crowding and inducing some passengers to depart earlier within the same discounted-fare period. It also changes crowding after the cutoff, affecting passengers whose fares remain unchanged. These responses produce a crowding profile that differs from the step change in fares. The first-stage estimates in Appendix Table C.5, column (2), show that the policy interactions predict expected crowding conditional on fare, rescheduling, and OD-by-date and date-by-time-bin fixed effects.

Using these interactions as instruments requires that, conditional on the model’s controls, they are unrelated to unobserved demand ξ_{jdt} . We use them alongside expected ICT in the multiple-instrument specification and in the random-coefficient model to help identify the income-dependent preference coefficients.

4.3 Sample and Summary Statistics

Our sample includes trips originating from the 16 EBD stations and 16 control stations and ending at any of the 342 stations in operation in 2016. We restrict to trips with a tap-in time between 6:30 and 8:30 AM. The EBD discount applies to passengers who tap in before 7 AM, so the two-hour window spans across the time cutoff. While the first train departs the terminal station around 5 AM on all four lines, ridership is low in the early hours (Figure 1), and we choose the earliest time of the sample at 6:30 AM. The 32-station restriction applies to the focal trips used in demand estimation. We measure crowding and construct ICPs using trips throughout the subway network, subject to the restrictions in Section 4.2.

The sample covers five non-consecutive collection periods spanning late February to December 2016. We exclude non-workdays because the EBD does not apply on those dates. These exclusions include June 9–10 for the Dragon Boat Festival and September 15–16 for the Mid-Autumn Festival, leaving 21 working days in the estimation sample.

Our baseline regressions study all riders departing from the selected stations during the morning peak. As a robustness check, we restrict the demand-estimation sample to frequent riders while continuing to construct the instrument from infrequent riders, as defined in Section 4.2. Results are reported in Appendix C. Crowding is always measured using the full universe of rides.

Each market is an OD pair on a date, with eight products corresponding to 15-minute departure-time bins. We retain markets with positive ridership in all eight bins. The final sample contains approximately 3.8 million trips across 40,191 markets. The EBD subsample includes 17,777 markets covering 1,621 unique OD pairs; the control subsample includes 22,414 markets covering 1,920 unique OD pairs. Regressions are weighted by the number of trips in each market.

Table 1: Summary Statistics

Variable	EBD Stations		Control Stations	
	Mean	SD	Mean	SD
# OD pairs	1,621		1,920	
# markets	17,777		22,414	
# passengers	10.96	21.81	12.41	16.28
# frequent passengers	10.23	20.96	11.57	15.62
Distance (km)	23.85	8.67	21.87	7.87
Total travel time (min)	56.33	19.80	47.81	16.16
In-train travel time (min)	39.67	15.18	37.31	14.34
Price (yuan)	5.21	1.15	5.45	0.84
Rescheduling, signed (min)	-0.28	11.01	-0.19	3.61
Rescheduling, absolute (min)	7.81	7.76	2.40	2.70
Crowding K (person·min/m ²)	132.75	65.32	123.25	52.07
Crowding density k (person/m ²)	3.31	1.30	3.44	1.28
Crowding density from ICP (infrequent)	0.08	0.07	0.08	0.08
Crowding density from ICP (total)	0.85	0.79	0.94	0.98
Crowding experience from ICT (infrequent)	1.48	1.26	1.30	1.17
Crowding experience from ICT (total)	14.97	13.99	14.62	14.51

Notes: Each observation is a product (jdt), defined as an OD pair–date–15-minute time bin. There are 142,216 observations in EBD covered OD pairs and 179,312 in control OD pairs. Means and standard deviations are unweighted across products. Crowding and companion measures describe realized exposure; demand estimation uses their expected-value proxies. ICP denotes incidental companion passengers; crowding density from ICP has a unit of persons/m². ICT stands for incidental companion time; crowding experience from ICT has a unit of person·min/m². “Infrequent” restricts to passengers with no more than three trips in the current week.

Table 1 reports product-level summary statistics for the key variables in the treatment and control groups. The two groups are broadly comparable. The average number of subway passengers per product is 11 in the EBD sample and 12.4 in the control sample. Trips from both groups are long, averaging 23.9 km and 56.3 minutes for those from EBD stations and 21.9 km and 47.8 minutes for those from control stations. This is not surprising because both EBD and control stations are at the far ends of suburban lines. The average in-train time is 39.7 minutes for trips from EBD stations and 37.3 minutes for those from control stations. The average fare is 5.21 RMB in the EBD sample and 5.45 RMB in the control sample, reflecting both the distance-based pricing schedule and the EBD

discount. The signed rescheduling deviation averages -0.28 minutes in the EBD sample (SD 11.01) and -0.19 minutes in the control sample (SD 3.61). The corresponding absolute deviations average 7.81 and 2.40 minutes (SDs 7.76 and 2.70), respectively.¹⁸ The negative mean rescheduling among the EBD sample is consistent with riders advancing their departures to qualify for the discount.

Passengers in the sample experience substantial crowding. The unweighted product-level average in-train crowding density is 3.31 persons per square meter for trips originating from EBD stations and 3.44 for those from control stations. The total in-train crowding experience, which integrates real-time crowding density over the duration of the trip (Figure 3), averages 132.8 person-minutes per square meter for EBD stations and 123.3 for control stations.¹⁹

ICPs, here defined as those with an origin O' upstream of O , and a destination D' at least eight kilometers away by straight-line distance, contribute meaningfully to crowding. For trips from EBD stations, total ICPs contribute a crowding density of 0.85 persons per square meter, accounting for roughly 26% of measured crowding. Infrequent riders, our preferred set of ICPs, contribute 0.08 persons per square meter. When weighted by the duration of shared travel, the incidental companion time of infrequent riders averages 1.48 person-minutes per square meter. The control sample has comparable figures.

5 Stylized Facts

This section documents several empirical patterns that motivate our modeling and identification choices. We show that the EBD policy effectively creates bunching in departure times and affects crowding, that crowding is predictable from regular commuting patterns, and that in-train crowding does not cause meaningful delays or rerouting. These facts support treating in-train crowding as a pure disamenity and validate our crowding measures, which assumes that passengers take the fastest route.

5.1 The EBD Creates Bunching in Departure Time and Crowding

Figure 1 plots the number of passengers entering, exiting, transferring through, and riding in trains at EBD and control stations by 5-min time bin during morning hours. At EBD stations (panel a), the entry flow displays a noticeable discontinuity around 7 AM. A cluster of passengers tap in just before the cutoff to qualify for the 30% discount, creating a visible dent in entries immediately afterward. This bunching in departure times in turn disturbs the distribution of passengers passing through the station around 7 AM. Because passengers passing through the station boarded the

¹⁸The model uses absolute deviations based on expected travel time $\widehat{TT}_{j,\text{dow},t}$. These can be nonzero for control riders even when their imputed optimal arrival time equals their actual arrival time.

¹⁹Integrated crowding is computed for each trip before aggregation to products. The average of integrated crowding does not generally equal the product of the separately averaged density and duration.

train at upstream stations and are themselves responding to the EBD cutoff, the resulting dent in the in-train crowding density (Appendix Figure C.1) is less sharp than the discontinuity in entry flows. The diffusion across upstream stations smooths the effect, but it also means that the EBD generates changes in crowding that ripple beyond the immediate neighborhood of the cutoff time. A passenger scheduled for an 8 AM departure who has no interest in taking the EBD may still adjust their departure time because the crowding on their regular trip is altered by other passengers responding to the EBD.

In contrast, panel (b) of Figure 1 shows that control stations exhibit no such discontinuity. The entry flow rises smoothly through the morning peak with no bunching around 7 AM, confirming that the EBD is the source of the departure-time distortion observed at treated stations.

5.2 Crowding Is Predictable

A key feature of the demand model is that passengers choose departure times based on *expected* crowding. Table 2 shows that most variation in crowding is captured by OD-by-time-bin fixed effects, with little additional explanatory power from conditioning on day of week. Crowding therefore varies little across observed dates for a given OD pair and departure-time bin. This regularity supports the assumption that experienced commuters can anticipate crowding. As described in Section 3.1, we proxy expected crowding by its average within each OD-by-day-of-week-by-time-bin cell.

These descriptive regressions use different fixed effects from the demand model. The demand model includes OD-by-date and date-by-time-bin fixed effects, which absorb daily demand differences across OD pairs and departure-time patterns common across OD pairs on each date. They leave differences in departure-time crowding profiles across OD pairs. We use expected ICT to instrument for expected crowding conditional on these fixed effects, fare, and rescheduling. The corresponding first-stage estimates are reported in Appendix Table C.5.

5.3 Crowding Does Not Slow Down the Train

Excessive time cost is a potential consequence of crowding and a confounding factor in the WTP to avoid crowding. In many transit settings, such as bus networks, high crowding coincides with road congestion, making it difficult to separate disutility from crowding per se from disutility caused by longer travel times. A distinctive feature of the current empirical setting is that in-train crowding does not meaningfully affect train speed, as trains run on dedicated tracks with fixed schedules, and delay is rare. This makes the Beijing Subway an ideal setting for identifying the pure disamenity cost of crowding. We test this premise directly.

We regress individual trip travel time on average in-train crowding density and in-station passenger counts, controlling for OD-by-date and line-time fixed effects. A one-person-per-square-meter increase in in-train crowding density is associated with only 0.27 additional minutes of travel time,

Table 2: Predictability of Crowding and Ridership

	Crowding Experience (K)			Crowding Density (k)		
	(1)	(2)	(3)	(4)	(5)	(6)
<i>Panel A: Crowding</i>						
R^2	0.8301	0.9876	0.9916	0.6511	0.9812	0.9870
Obs. (million)	5.79	5.74	4.96	5.79	5.74	4.96
	Frequent Riders			Infrequent Riders		
	(1)	(2)	(3)	(4)	(5)	(6)
<i>Panel B: Ridership</i>						
R^2	0.6804	0.9342	0.9449	0.4317	0.5226	0.6281
Obs. (million)	12.53	12.53	11.93	12.53	12.53	11.93
OD FE	X			X		
OD-Time FE		X			X	
OD-DOW-Time FE			X			X

Notes: Panel A reports R^2 from regressions of realized crowding on fixed effects of increasing granularity. Crowding experience K is measured in person-min/m² and crowding density k in persons/m². Panel B reports analogous regressions for ridership, separately for frequent riders (≥ 4 trips in the current week) and infrequent riders.

relative to an average of about 39 minutes (Appendix Table C.6, column 1). To put the magnitude in perspective, we convert the delay cost into monetary terms. At an assumed value of time of 0.4 RMB per minute, the implied monetary cost is about 0.11 RMB, an order of magnitude smaller than the willingness to pay for reducing discomfort from crowding. Appendix C.2 reports the calculation and additional evidence. Although we do not directly estimate the disutility from in-station crowding, we show in the robustness checks that in-station crowding does not confound the estimated WTP to avoid crowded trains.

5.4 Crowding Does Not Result in Rerouting

The computation of in-train crowding relies on a train assignment algorithm that assigns each passenger to the fastest route. A natural concern is that some passengers may deviate from this route in response to crowding. If passengers systematically choose longer but less crowded alternatives, the measured crowding on their assigned route would overstate the crowding they actually experience, and the degree of overstatement would increase with crowding levels.

The distributions of travel time are tightly concentrated around the OD-specific mean and remain similar across crowding quartiles (Appendix Figures C.9 and C.10). A one-person-per-square-meter increase in crowding density is associated with a 0.12-minute increase in absolute deviation from the OD-date-time-bin cell mean, relative to a mean absolute deviation of 2.40 minutes (Appendix

Table C.7). These patterns support treating substantial crowding-induced rerouting as uncommon. Appendix C.3 presents the detailed evidence.

6 Estimation Results

6.1 Constant Coefficient Model Estimates

Table 3 reports estimates of the constant coefficient logit model in Equation (4). Crowding and ICT in this section refer to expected values, proxied by the OD-by-day-of-week-by-departure-time-bin averages defined in Sections 3.1 and 4.2; table row labels omit “expected”. The first instrument set uses ICT alone; the second adds the four EBD policy interactions described in Section 4.2. Both specifications control for fare, rescheduling, OD-by-date fixed effects, and date-by-time-bin fixed effects.

Column (1) presents uninstrumented OLS estimates with OD-by-date and date-by-time-bin fixed effects. The price coefficient is negative and precisely estimated, indicating that passengers are less likely to choose a time bin with a higher fare. The rescheduling coefficient is also negative and statistically significant at conventional levels. The crowding coefficient, however, is *positive*, and the implied MWTP for crowding reduction is correspondingly negative. Rather than reflecting a genuine preference for crowded trains, the “wrong” sign highlights the endogeneity of crowding: trips taken at crowded times and on crowded routes carry high unobserved value. The positive correlation between crowding and unobserved demand confounds the true disutility, biasing the OLS estimate of β upward.

We use expected ICT to instrument for expected crowding, as described in Section 4.2. Our preferred ICT measure uses infrequent ICPs who board upstream and travel to destinations at least 8 km from the focal destination. Section 6.2 examines sensitivity to these restrictions.

The remaining columns of Table 3 report GMM estimates. Columns (2) and (3) use ICT alone and report the first-stage and second-stage results, respectively. Columns (4) and (5) add the four EBD policy instruments and again report the first and second stages. Column (6) uses the same instruments as columns (4) and (5), replacing the linear crowding term with its square.

In the first stage (columns 2 and 4), ICT is a strong positive predictor of crowding, with coefficients of 11.3 and 10.9 and first-stage Wald F-statistics of 607 and 299 in the ICT-only and ICT-plus-EBD specifications, respectively. The F-statistics comfortably exceed conventional thresholds for weak instruments. The price coefficient in the ICT-plus-EBD first stage (column 4) is much larger in magnitude than in the ICT-only first stage (column 2) because the additional instruments, the EBD policy interactions with time trends, are highly collinear with price. The EBD discount is itself a function of treatment status and time of day, so the policy interactions absorb most of the price variation, leaving the first-stage price coefficient identified off a smaller and more idiosyncratic residual.

Table 3: Constant Coefficient Model Estimates

	OLS: Uninstrumented	GMM: ICT only		GMM: ICT + EBD instruments		GMM: Nonlinear
	(1)	1st Stage (2)	2nd Stage (3)	1st Stage (4)	2nd Stage (5)	2nd Stage (6)
Crowding	0.0058 (0.0003)		-0.0191 (0.0018)		-0.0092 (0.0011)	
Crowding ²						-0.000028 (0.000002)
Price	-0.2046 (0.0141)	-5.0535 (0.4153)	-0.3146 (0.0182)	29.2149 (2.0870)	-0.1754 (0.0087)	-0.1457 (0.0084)
Rescheduling	-0.0297 (0.0009)	0.0229 (0.0278)	-0.0278 (0.0012)	-0.1363 (0.0221)	-0.0253 (0.0009)	-0.0256 (0.0009)
ICT		11.2860 (0.4580)		10.9076 (0.4441)		
Other IVs				✓		
First-stage F -statistic		607.1		299.4		479.3
Overall R^2	0.659		0.523		0.608	0.618
OD-Date FE	X	X	X	X	X	X
Date-Time FE	X	X	X	X	X	X
Mean Rescheduling (mins)	4.796		4.796		4.796	4.796
Mean Crowding (person-min/m ²)	127.4		127.4		127.4	127.4
<i>Implied MWTP</i>						
Crowding	-0.0284 (0.00250)		0.0607 (0.00540)		0.0526 (0.00597)	<i>see</i> <i>Fig. C.7</i>
Rescheduling	0.145 (0.0111)		0.0882 (0.00678)		0.144 (0.00883)	0.1755 (0.0120)

Notes: This table reports OLS and GMM estimates of Equation (4) for the overall sample of riders. Column (1) reports uninstrumented OLS estimates. In columns (2)–(3), crowding is instrumented by ICT from infrequent riders with $|D - D'| \geq 8$ km. In columns (4)–(6), four EBD policy interactions are added as instruments ($\text{Treat} \times 1\{t < 7:00\}$, $\text{Treat} \times \text{trend}$, $\text{Treat} \times \text{trend}^2$, $\text{Treat} \times \text{trend}^3$). Column (6) includes a quadratic crowding term. Overall R^2 includes the fixed effects and is computed as one minus the ratio of the structural residual sum of squares to the total centered sum of squares of log shares. The F -statistics assess first-stage instrument relevance; column (6) reports the statistic for the quadratic specification. Standard errors in parentheses, clustered at the OD level. MWTP computed as the ratio of the relevant coefficient to the price coefficient. Appendix Table C.2 reports a parallel set of estimates restricted to the sub-sample of frequent riders. Crowding and ICT denote expected values; the row labels omit "expected".

Turning to the second stage, the crowding coefficient is negative in all GMM specifications, in sharp contrast to the positive OLS estimate in column (1). With ICT alone (column 3), the crowding coefficient is -0.0191 and the price coefficient is -0.315, yielding an implied MWTP for crowding reduction of 0.061 RMB per person-min/m². The implied MWTP to reduce rescheduling by one minute is 0.0882 RMB. The ICT-plus-EBD specification (column 5) produces a similar implied MWTP of 0.053 RMB for crowding reduction and 0.144 RMB for rescheduling reduction. The ICT-plus-EBD specification is our preferred specification.

To gauge the economic magnitudes, consider a representative trip lasting approximately 38 minutes in train, using the unweighted product-level averages in Table 1. The willingness to pay to reduce crowding density by one person per square meter over the course of the trip is 0.053×38

≈ 2.0 RMB, or roughly 38% of the average fare of 5.3 RMB. Assuming linear marginal disutility, the total crowding disutility for a trip at the mean total crowding of about 127 person·min/m² amounts to $0.053 \times 127 \approx 6.7$ RMB, exceeding the fare itself. These magnitudes confirm that in-train crowding is a first-order determinant of the value of a subway trip. For rescheduling, deviating from the optimal arrival time by 15 minutes costs $0.144 \times 15 \approx 2.2$ RMB. Relative to the median hourly wage of about 43 RMB (7,500 RMB per month over 22 eight-hour workdays), the rescheduling cost of a 15-minute deviation amounts to about 5% of hourly earnings. The relatively small rescheduling cost may be due to the fact that most rescheduling in our setting is departing and arriving earlier than optimal. Departing in the neighborhood of 7 AM, virtually all trips conclude before 9 AM, the start of conventional work hours. Rescheduling to a later arrival time after the start of conventional work hours may incur a larger cost. Nevertheless, those findings suggest that passengers face a meaningful tradeoff between departing at a less crowded time and arriving closer to their preferred schedule.

Column (6) allows for a non-linear effect of crowding by including the quadratic term of the crowding experience. In this specification, we omit the linear crowding term and include the quadratic term only. This essentially assumes that the marginal disutility is zero when crowding density is zero, which is reasonable. The estimated coefficient on the quadratic term is negative and precisely estimated, indicating that crowding disutility is convex and its marginal cost increases with exposure. The MWTP for crowding reduction varies with the level of crowding, as shown in Figure C.7 in Appendix C. At the 25th percentile of expected cumulative crowding exposure, the MWTP is 0.032 RMB per person·min/m², rising to 0.049 at the median and 0.086 at the 95th percentile. The increasing pattern implies that the welfare cost of crowding is concentrated on the most congested trips.


Our estimates of the crowding cost are broadly consistent with the transportation literature, although the comparison requires care, since most prior work expresses crowding as a multiplier on the value of in-vehicle time rather than as a direct disutility per unit of density. Stated preference studies find that this multiplier roughly doubles once passengers stand in dense conditions (Wardman and Whelan, 2011) and that the cost is convex in density (Haywood and Koning, 2015), mirroring the increasing marginal disutility in our quadratic specification, and revealed preference studies using smartcard data reach similar conclusions (Hörcher et al., 2017; Yap et al., 2020). Using the assumed value of uncrowded travel time of 0.4 RMB per minute from Appendix C.2, the preferred coefficient implies a crowding multiplier of about 1.45 at the mean density of 3.4 persons per square meter and 1.66 at the ninetieth percentile of about 5 persons per square meter.

6.2 Robustness

Alternative ICP Definitions. We examine whether the crowding MWTP estimates are sensitive to the definition of ICPs. The potential endogeneity arises from demand shocks correlated at the

origin, the destination, or through common travel purpose. To examine these sources, we vary the construction of the instrument by progressively tightening restrictions on which co-riders qualify as ICPs.

Figure 6: Crowding MWTP across ICP Restrictions



Notes: Each point reports the estimated MWTP for reducing crowding by one person per square meter for one minute (RMB), from a separate regression on the overall demand sample. Both panels restrict ICPs to upstream origins. The horizontal axis varies the minimum straight-line distance $|D - D'|$ between focal and ICP destinations. Blue dots report uninstrumented OLS estimates; red dots use expected ICT from all qualifying ICPs; green dots use expected ICT from infrequent ICPs (≤ 3 trips in the current week). Both GMM series include the four EBD policy instruments. Vertical bars show 95% confidence intervals. Standard errors are clustered at the OD level.

Figure 6 compares crowding MWTP estimates across restrictions on the passengers used to construct ICT. Both panels use ICPs who board upstream of the focal origin O . The horizontal axis increases the minimum straight-line distance between the focal destination D and the ICP destination D' ; an 8-km threshold excludes ICPs whose destinations are less than 8 km from D . Panel (b) additionally excludes ICP origins within 2 km of O . The red estimates use all qualifying ICPs, while the green estimates restrict ICPs to infrequent riders. The demand estimation sample remains the overall sample in both cases. All GMM estimates use ICT plus the four EBD policy instruments. The blue dots report the uninstrumented OLS benchmark.

When we only require $D \neq D'$ but distance $|D - D'|$ is not restricted (the leftmost point), the GMM estimate with ICT as IV yields a negative MWTP, with a magnitude similar to the OLS estimate. Using only infrequent ICPs pushes the estimate upward, but the MWTP remains negative and statistically significant. This suggests that ICPs who travel to destinations close to D are likely responding to correlated demand shocks. Restricting to infrequent ICPs improves, but does not fully eliminate the endogeneity concern.

As we gradually remove ICPs whose destinations are close to D , the GMM estimate rises steadily. By restricting $|D - D'|$ to exceed 3 km, the MWTP for crowding reduction becomes positive. Restricting $|D - D'|$ to be greater than approximately 7 to 8 km, the estimates converge. With $|D - D'|$ sufficiently large, estimates using either all ICPs or infrequent ICPs as instruments also converge.

This is intuitive. Even when ICPs are traveling for the same purpose (say arriving at work on time), the fact that they are going to a faraway destination means that they are responding to an unrelated demand shock. Thus further restricting to infrequent ICPs does not make a further difference. When $|D - D'|$ is greater than 8 km, the MWTP stabilizes at roughly 0.05 RMB per person·min/m². Further tightening the distance restriction does not meaningfully change the estimates.

Panel (b) repeats the exercise while further requiring ICPs to originate from a station O' more than 2 km away from the focal passenger's origin O . The results are nearly identical to those in panel (a), suggesting that the source of endogeneity is not from correlated demand shocks at the origin. Taken together, the two panels show that the endogeneity of crowding stems primarily from common shocks at the *destination*. Intuitively, the most important shock that drives crowded subway cars during morning rush hours is the need to arrive at work at a required time, which concentrates travel demand among passengers heading to nearby destinations during the same narrow window.

These convergence patterns support the preferred 8-km destination restriction used in Table 3. This restriction creates an exclusion zone of roughly 200 square kilometers around the focal destination. One may worry that it leaves too few qualifying ICPs to generate a strong first stage. This concern is mitigated by the scale of the Beijing Subway network. Figure C.6 in Appendix C shows that the EBD stations are at the far suburban ends of the network, while the vast majority of the system extends well beyond the 8-km radius from the city center (panel a). Because destinations are spread throughout this extensive network, the 8-km exclusion zone around each focal destination D removes only a small share of potential ICPs. Indeed, the median infrequent ICP alights at a station roughly 11 km from D (panel b), well beyond the exclusion threshold. As a result, the first-stage relationship remains strong, with Wald F-statistics exceeding 290 in all specifications (Table 3).

ICP as Instrument. Our preferred instrument uses incidental companion *time* (ICT), which weights each ICP by the duration of shared travel. A natural alternative is to use the expected count of ICPs, discarding the time weights. Table C.3 in Appendix C reports both ICP-only and ICP-plus-EBD specifications using ICP. In the first stage, ICP is a strong predictor of crowding, with a coefficient of approximately 179 and Wald F -statistics of 488.7 (ICP-only) and 302.9 (ICP-plus-EBD). The MWTP for crowding reduction is 0.057 RMB per person·min/m² in the ICP-only specification and 0.048 under ICP-plus-EBD. These estimates are close to their ICT counterparts (0.061 and 0.053, respectively), as expected because ICP and ICT exploit the same source of exogenous variation. The stability of the crowding WTP across the two instrument constructions is reassuring.

Frequent riders. The baseline demand model is estimated on trips from all users, while the instrument is constructed from the expected incidental companion time of infrequent riders alone. As a robustness check, we estimate the demand model on frequent riders while continuing to construct the instrument from infrequent riders. This checks whether the estimates depend on including in-

frequent riders in the demand sample.

Appendix Table C.2 reports the parallel set of estimates. The implied MWTP for crowding reduction is 0.061 RMB per person·min/m² with ICT alone and 0.052 with ICT plus the EBD instruments, essentially identical to the overall-sample values of 0.061 and 0.053 in Table 3. We retain the overall demand sample for the remainder of the paper.

Controlling for In-Station Crowding. A potential concern is that in-train crowding may be correlated with in-station crowding, which could independently affect route choice. In-station crowding is qualitatively different from in-train crowding: passengers experience crowded stations in areas where the per-capita space is substantially larger than inside a train car. The main channel through which station congestion could plausibly affect riders is via delays in gates, corridors, and platforms; Section 5.3 shows directly that any such delay is small in magnitude in our setting (a 35 percent increase in in-station passenger counts is associated with about 0.80 additional minutes of travel time; Appendix C.2). We use a unified definition of in-station crowding: InStn_{jdt} is the total number of passengers that the focal passenger choosing product jdt encounters in stations during entry, transfer, and exit, weighted by the duration of the overlap, and has a unit of thousand persons-minutes.

Table C.4 reports two specifications that add in-station crowding to the estimating equation for the frequent-rider sample. Column (1) adds InStn_{jdt} as an uninstrumented control under the ICT-only specification, while Column (2) treats InStn_{jdt} as a second endogenous regressor, jointly instrumented with in-train crowding by the multiple IVs we have. The EBD policy instruments shift entry timing, which moves both in-train and in-station crowding through the same exogenous mechanism, providing the variation needed to identify both endogenous regressors. The Wald F-statistic is 470.6. The estimated MWTP to reduce in-train crowding remains positive after controlling for in-station crowding. The MWTP to reduce crowding by one person·min/m² is 0.05 RMB in Column (1) and 0.0482 RMB in Column (2), compared with 0.061 and 0.052 in the corresponding frequent-rider baseline specifications.

In Column (1), the in-station crowding coefficient is positive (+0.040). A positive coefficient on a disamenity implies the presence of endogeneity: high-demand markets attract more riders, leading to crowded stations and biasing the InStn_{jdt} estimate upward. In Column (2), the in-station coefficient turns negative, and the implied MWTP for in-station crowding reduction is about 0.048 RMB per thousand person·min reduction. Evaluated at the robustness-sample mean (14.3 thousand person·min), the implied disutility from average in-station crowding is roughly $0.048 \times 14.3 \approx 0.69$ RMB, an order of magnitude smaller than the average in-train disutility computed in Section 6.1.

6.3 Heterogeneous Willingness to Pay

The standard logit model restricts all passengers to share the same preference parameters. We now relax this assumption and allow the marginal utility of price, crowding, and rescheduling to vary with income, as specified in the random coefficient model of Equation (5). Each preference coefficient is specified as a linear function of simulated passenger income, recentered at the sample mean of 8.77 thousand RMB. The model is estimated by GMM with the same ICT-plus-EBD instrument set used in the preferred specification, following the BLP contraction mapping procedure (Berry et al., 1995).

Table 4 reports the estimation results. The price intercept $\hat{\alpha}_0 = -0.179$ indicates that the average passenger derives disutility from fare, while the positive income interaction $\hat{\alpha}_1 = 0.013$ implies that higher-income riders are less price sensitive. The crowding intercept $\hat{\beta}_0 = -0.026$ indicates that crowding reduces utility, and the negative income interaction $\hat{\beta}_1 = -0.009$ reveals that higher-income passengers experience substantially stronger crowding disutility. The rescheduling coefficient $\hat{\rho}_0 = -0.0176$ is negative, but the income interaction $\hat{\rho}_1 = 0.0015$ is not statistically different from zero, suggesting that the cost of deviating from the optimal arrival time is relatively uniform across income levels.


Table 4: Random Coefficient Model Estimates

	(1)
Price (α_0)	-0.1790 (0.0109)
Price $\times$ Income (α_1)	+0.0130 (0.0047)
Crowding (β_0)	-0.0256 (0.0006)
Crowding $\times$ Income (β_1)	-0.0093 (0.0002)
Rescheduling (ρ_0)	-0.0176 (0.0058)
Rescheduling $\times$ Income (ρ_1)	+0.0015 (0.0019)
GMM Objective	0.3030
OD-Date FE	X
Date-Time FE	X

Notes: This table reports the random coefficient model estimates of Equation (5). Income is measured in thousands of RMB per month and recentered at the sample mean (8.77 thousand RMB). The estimates use the overall sample with OD-by-date and date-by-time-bin fixed effects. Heteroskedasticity-robust standard errors in parentheses. The model is estimated using the ICT-plus-EBD instrument set (ICT from infrequent riders with $|D - D'| \geq 8$ km, plus four EBD policy interactions). Crowding and ICT denote expected values; the row labels omit "expected".

Figure 7 plots the estimated MWTP for crowding reduction and rescheduling avoidance as a

Figure 7: Willingness to Pay by Income



Notes: The solid line shows the point estimate of MWTP as a function of monthly income. Dashed lines indicate the 95% confidence interval (delta method). Vertical dashed lines mark selected percentiles of the income distribution. Panel (a) reports the MWTP for reducing in-train crowding by one person-min/m² (in RMB). Panel (b) reports the MWTP for reducing schedule deviation by one minute (in RMB). The plots are truncated at 15 thousand RMB (between the 75th and 90th percentiles), beyond which the price coefficient $\alpha(I)$ approaches zero and the WTP ratio becomes imprecisely estimated.

function of monthly income, with 95 percent confidence intervals computed by the delta method. The MWTP for crowding reduction is steeply increasing in income (panel a). At the median income of 7.5 thousand RMB, the MWTP is approximately 0.07 RMB per person-min/m², close to the estimate in the homogeneous setting of 0.053 from Table 3. At the 90th percentile of income (17.5 thousand RMB), the MWTP rises to roughly 1.62 RMB per person-min/m², more than twenty times the median value.²⁰

In contrast, the MWTP for rescheduling avoidance is relatively flat across income levels (panel b), hovering around 0.10 RMB per minute throughout most of the income distribution. The stability of the rescheduling WTP is consistent with the insignificant income interaction $\hat{\rho}_1$ and suggests that the cost of arriving at an inconvenient time is not strongly differentiated by income.

The steep income gradient in crowding WTP has important implications for policy. It suggests that higher-income passengers would be willing to pay a substantial premium for less crowded service, while lower-income passengers are relatively indifferent to crowding. This pattern provides a potential rationale for service differentiation, where higher-income riders self-select into a more expensive but less crowded option. We explore this possibility in the welfare analysis of Section 7.

²⁰The graph also shows that the MWTP to reduce crowding turns negative for passengers with monthly income under 6,000 RMB, or about 40th percentile of the income distribution. The substantial share of users with negative MWTP and the steep gradient with regard to income are mostly due to the underlying income data, which is imputed by Baidu Maps with limited and indirect information. In welfare calculation and counterfactual analysis below, we define high- and low-income passengers from income percentiles to avoid negative or very high values of MWTP.

7 Policy Counterfactuals

Subway crowding is a congestion externality. Each additional rider imposes discomfort on fellow passengers that is not incorporated in the fare. In this section, we compare three policies that address this externality: a Pigouvian crowding tax, capacity rationing through queuing, and a two-class train configuration with differentiated fares.

We evaluate these policies using a simplified version of the estimated demand model. We replace the subway network with a single origin–destination pair and collapse the morning peak into one representative departure period. Passengers choose between the subway and an outside option, with ridership and crowding determined in equilibrium. The model retains differences in fare sensitivity and crowding valuations across income groups but abstracts from departure-time adjustments and interactions across routes. These simplifications avoid solving simultaneously for passenger choices and crowding throughout the network and across departure periods.

The simplified model preserves the mechanisms underlying the policy comparison. A crowding tax makes passengers pay for the discomfort they impose on others. Queuing limits ridership through waiting and schedule-disruption costs. Differentiated fares and car densities allow passengers to sort according to their crowding valuations. The quantitative effects nevertheless depend on the margins available to passengers. In the full setting, passengers may respond by traveling at another time or taking another route; here, they can leave the subway or, under the two-class policy, switch classes. The exercise therefore illustrates the aggregate and distributional effects of these policy mechanisms, rather than determining optimal fares across the network and over time.

7.1 Setup

Consider a stylized linear city where all residents live in location O and work in location D . O and D are connected by a subway line and a surface road. There are two income groups $g \in \{\ell, h\}$ that equally divide the entire population, $M_\ell = M_h = M/2$, where M denotes the total market size. Each resident commutes to work. They either ride the subway or take the outside option, the utility from which is standardized to one.

Deterministic utility from choosing the subway is

$$V_g = A + \alpha_g P + \underbrace{\beta_g \frac{\text{TravelTime}}{\text{CarArea} \times J}}_L N, \quad (11)$$

where A is a common mean utility, P is the fare, N is the total subway ridership on the segment, which is an endogenous outcome of the equilibrium. Let $L = \text{TravelTime}/(\text{CarArea} \times J)$, where J is the number of cars in a train, thus the denominator indicates the total floor space in a train. NL measures the total crowding exposure a passenger experiences during the ride (crowding density

times minutes). In this simple static setting, commuters travel at the same time, so departure from optimal arrival time does not enter the utility. We relax this assumption when we introduce queuing in the quantity control policy experiment. Passengers derive disutility from price and crowding, $\alpha_g < 0$ and $\beta_g < 0$. Idiosyncratic taste shocks are i.i.d. type-I extreme value, so the logit share for group g is $s_g(V_g) = e^{V_g} / (1 + e^{V_g})$ and group ridership is $N_g = M_g s_g(V_g)$, with $N = \sum_g N_g$.

We split riders at the 41st and 59th percentiles of the simulated income distribution from the Baidu commuting data.²¹ The low-income group has a monthly income of $y_\ell = 6.19$ thousand RMB, while the high-income group has $y_h = 9.11$ thousand RMB. The BLP estimates in Table 4 imply $(\alpha_\ell, \beta_\ell) = (-0.213, -0.002)$ for the low-income group and $(\alpha_h, \beta_h) = (-0.174, -0.029)$ for the high-income group.

Baseline ridership follows directly from the data. The ridership-weighted average crowding density is 4.03 persons per square meter, car area is 53.2 m², and a train has $J = 6$ cars, yielding an equilibrium ridership of $N_{eq} = 4.03 \times 53.2 \times 6 = 1,286$. Given a ridership-weighted subway mode share of 0.43 (constructed from the BHTS imputation on our working sample), the market size is $M = 1,286 / 0.43 \approx 2,983$. Ridership-weighted average in-train travel time of 31.83 minutes gives $L = 31.83 / (53.2 \times 6) = 0.1$. The ridership-weighted baseline fare is $P_{base} = 5.08$ RMB, which we assume is the marginal operating cost. The mean utility A is calibrated so that the private equilibrium at P_{base} reproduces N_{eq} , which yields $A = 1.77$. Table 5 reports parameter values used in the welfare and counterfactual exercises.

For the two-class counterfactual we partition the six-car train into 3 Business and 3 Standard cars. The symmetric split is a natural benchmark because $L_B = L_S = 2L$ and the two classes are identical apart from the fare. Riding in either class has the same mean utility A . Class-specific utilities are $V_{gB} = A + \alpha_g P_B + \beta_g L_B N_B$ and $V_{gS} = A + \alpha_g P_S + \beta_g L_S N_S$. Any crowding difference between the two classes is therefore driven entirely by self-selection in response to the fare gap, not by a mechanical capacity asymmetry or differences in other features.

Through the well-known independence of irrelevant alternatives (IIA) property, a standard multinomial logit would mechanically inflate the subway share once the two-class configuration is introduced. This is because even if prices are identical, draws of idiosyncratic preferences over the two subway classes draw passengers into each subway class. We therefore use a nested logit, where the two classes of subway cars form the lower nest, while subway and other modes form the upper nest. The inclusive value for group g is

$$IV_g = \lambda \ln(e^{V_{gB}/\lambda} + e^{V_{gS}/\lambda}), \quad (12)$$

and the corresponding choice probabilities follow the standard nested-logit form with nesting pa-

²¹The split p41/p59 is symmetric around the median (41 + 59 = 100). The 41st percentile is chosen because it sits just above the point at which the crowding coefficient changes sign, so that low-income riders have a weakly non-negative crowding coefficient while high-income riders strictly dislike crowding. Because low-income riders defined in our case derive little disutility from crowding, and because the MWTP increases sharply

Table 5: Parameter Values Used in the Counterfactual Simulations

Parameter	Value	Description
<i>Structural parameters (from Table 4)</i>		
α_0	-0.179	Price coefficient, intercept
α_1	+0.013	Price coefficient, income slope
β_0	-0.026	Crowding coefficient, intercept
β_1	-0.009	Crowding coefficient, income slope
ρ_0	-0.017	Rescheduling coefficient, intercept
ρ_1	+0.002	Rescheduling coefficient, income slope
<i>Group-level parameters, evaluated at demeaned incomes</i>		
α_ℓ, α_h	-0.213, -0.174	Price sensitivity by group
β_ℓ, β_h	-0.002, -0.029	Crowding disutility by group
<i>Market and calibration</i>		
$y_\ell, y_h, \bar{y}$	6.19, 9.11, 8.77 (k RMB)	Group incomes (p41/p59) and sample mean
N_{eq}	1,286	Baseline equilibrium ridership
M	2,983	Total market size
L	0.100	Congestion exposure (min/m ² /person)
P_{base}	5.08 RMB	Baseline fare
A	1.77	Calibrated mean utility
<i>Two-class configuration</i>		
J_B, J_S	3, 3	Business, Standard car count
λ	0.5	Nesting parameter

Notes: Structural parameters are from the random-coefficient logit estimates in Table 4. Group-level coefficients are evaluated at demeaned incomes $y_g - \bar{y}$ where $\bar{y} = 8.77$ thousand RMB is the sample mean. Baseline ridership, market size, and congestion exposure are computed from the sample. The mean utility A is calibrated to reproduce N_{eq} at P_{base} .


parameter $\lambda \in (0, 1]$. When $\lambda = 1$, the model collapses to multinomial logit. When $\lambda < 1$, Business and Standard classes are more correlated with each other than either is with the outside option, which attenuates the spurious variety bonus. We use $\lambda = 0.5$ as our preferred value, at the lower end of the range typical for sub-modes within a transit nest in the discrete-choice transport literature, and examine sensitivity to $\lambda \in [0.3, 0.8]$ in the appendix.

7.2 Private Demand, Social Demand, and Externality

The welfare analysis builds on the distinction between marginal private demand and marginal social demand, illustrated in Figure 8. In the market equilibrium, the price equals each rider’s marginal willingness to pay net of the crowding disutility the rider experiences personally. Each additional passenger, however, also imposes crowding on every other rider, a cost the private choice does not

internalize. The social planner accounts for this externality, and the vertical distance between private demand and social demand captures the marginal externality.

Figure 8: Crowding Externality and the Optimal Crowding Tax



Notes: MPD denotes marginal private demand, MSD marginal social demand, and the vertical distance between them is the externality wedge. At the social optimum N^* , the Pigouvian crowding tax equals the wedge $w(N^*)$, and the deadweight loss from unpriiced crowding is the shaded triangle between MPD and MSD over $[N^*, N_{eq}]$.

To construct the inverse private demand, fix any candidate aggregate ridership N . The price that rationalizes N as a private equilibrium must satisfy condition $N = \sum_g M_g s_g(A + \alpha_g P + \beta_g L N)$, which implicitly defines the mapping $PD(N)$. Letting $N_g(P, N) = M_g s_g(A + \alpha_g P + \beta_g L N)$ denote group-specific ridership, the difference between private and social demand, thus the negative externality, can be written as

$$w(N) \equiv PD(N) - MSD(N) = L \sum_{g \in \{\ell, h\}} \frac{\beta_g}{\alpha_g} N_g(PD(N), N). \quad (13)$$

The intuition behind Equation (13) is straightforward. Each additional rider raises crowding linearly in L , and the monetized disutility per unit of crowding for group g is β_g/α_g . Summing these monetized effects across groups, weighted by the number of riders exposed, gives the aggregate marginal external cost.

7.3 Counterfactual Policies

We evaluate three policy instruments against the baseline: a Pigouvian crowding tax, capacity-rationing queuing, and a two-class pricing scheme.

We assume P_{base} is the operator's marginal cost, and the operator cannot operate at a loss. The constrained social optimum solves

$$PD(N^*) - w(N^*) = P_{\text{base}}, \quad (14)$$

and the corresponding Pigouvian crowding tax equals the externality at the socially optimal level of ridership, $T^* = w(N^*)$. Riders face $P_{\text{cons}} = P_{\text{base}} + T^*$. At P_{cons} , individual choices under the competitive equilibrium achieves the social optimal levels. Tax revenue is $T^* N^*$.

Capacity-rationing queuing is a non-price instrument in which each train has a fixed boarding capacity equal to N^* and excess demand produces a queue of length W in minutes. The wait enters utility through two channels: as a time cost priced at the group-specific value of time VOT_g , which we take as half of the wage rate, and as a schedule-disruption cost captured by the rescheduling coefficient ρ_g estimated in Section 6.3. The generalized price is $P_{\text{base}} + \text{VOT}_g \times W$, and utility includes an additional term $\rho_g \times W$. The equilibrium wait W^* clears the queue, $\sum_g M_g s_g(V_g(W^*)) = N^*$, and we solve by bisection on W . Queuing generates no transferable revenue, and the wait cost is an additional deadweight loss; users decide to join the queue with rational expectations about the equilibrium waiting time.

The two-class counterfactual maintains revenue neutrality $P_B N_B + P_S N_S = P_{\text{base}} N_{\text{eq}}$. For any given P_S , we find $P_B > P_S$ and solve the nested-logit fixed point in (N_B, N_S) , and retain the pair that maximizes the variety-adjusted consumer surplus defined in Section 7.4. The numerical implementation is described in Appendix D.

7.4 Welfare Calculation

We now describe how consumer surplus and total welfare are computed for the baseline and three counterfactuals introduced above. Because the idiosyncratic taste shocks are type-I extreme value, consumer surplus in each group admits the standard log-sum form. In the single-class case, which includes the baseline, the Pigouvian tax, and the queuing counterfactual,

$$CS_g(N, P) = -\frac{M_g}{\alpha_g} \ln(1 + e^{A + \alpha_g P + \beta_g L N}), \quad CS = \sum_g CS_g. \quad (15)$$

The factor $1/\alpha_g$ converts utility into monetary units, and the crowding externality is already embedded in V_g . In the two-class nested logit, the log-sum argument is replaced by the inclusive value, giving $CS_g = -(M_g/\alpha_g) \ln(1 + e^{IV_g})$.

The nested logit formula mechanically exceeds single-class consumer surplus even when $P_B = P_S$, because the inclusive value exceeds the deterministic utility according to Jensen’s inequality. This artificial variety bonus is a known feature of discrete-choice models, which tend to overstate the welfare gains from expanding the choice set when the additional alternatives do not correspond to genuine product differentiation. Because relabeling physically identical cars as Business or Standard does not constitute real variety, we remove this artifact. Let VB denote the consumer-surplus gain when we impose the revenue-neutral equal price, $P_B = P_S = P_{eq}$. The variety-adjusted consumer surplus is

$$CS_{VA} = CS_{total} - VB, \quad (16)$$

which by construction equals baseline consumer surplus in the absence of price differentiation. We report CS_{VA} as the primary welfare measure for two-class counterfactuals. Total welfare is reported as $W = CS + \text{TaxRev}$, with CS_{VA} substituting for CS in the two-class case. Under revenue neutrality, fare revenue is held fixed across counterfactuals.

7.5 External Road Congestion

Changes in subway ridership also affect surface-road congestion as displaced riders shift to other modes. Let $MECC_{veh-km}$ denote the marginal external cost of road congestion per vehicle-kilometer, $\bar{d}$ the mean trip distance in kilometers, and κ the average car occupancy, so that the passenger-level congestion cost in cars is $c_{car} = MECC_{veh-km}/\kappa$. Diversion shares to other modes differ across groups. We assume that low-income riders shift to $p_{bus,\ell} = 0.5$, $p_{car,\ell} = 0.25$, and other modes with probability 0.25, while high-income riders shift to $p_{bus,h} = 0.15$, $p_{car,h} = 0.75$, and other modes with probability 0.10, reflecting the much higher rate of car ownership among higher-income households. Following Gu et al. (2024), each bus per kilometer generates 0.4 car-equivalents of congestion, while riders shifting to other modes are assumed not to affect road traffic. The change in road-congestion cost relative to baseline is

$$\Delta MECC = c_{car} \bar{d} \sum_g (p_{car,g} + 0.4 p_{bus,g}) (N_g^{base} - N_g^{cf}). \quad (17)$$

We report welfare both without and with MECC, as W and $W - \Delta MECC$.

7.6 Counterfactual Results

Table 6 summarizes the comparison across policies. Panel A reports prices and ridership, and Panel B reports consumer surplus, revenue, and welfare.

The Pigouvian tax, at 3.3 RMB per trip, lowers ridership to the socially optimal level of $N^* = 1,077$, a 16 percent reduction from baseline. Consumer surplus falls from 10,885 to 7,671 RMB, but the tax

Table 6: Welfare Comparison across Counterfactual Scenarios

Panel A: Prices and Ridership

	P	tax	W^*	N	N_ℓ	N_h	k
<i>1. Single-class configuration</i>							
1.1 Baseline (status quo)	5.08	–	–	1,286	1,132	154	4.03
1.2 Optimal crowding tax	8.38	3.30	–	1,077	919	158	3.37
1.3 Quantity control (queuing)	5.08	–	7.75	1,077	938	139	3.37
<i>2. Two-class (3+3, $\lambda = 0.5$)</i>							
2.1 Welfare maximization	–	–	–	1,524	1,235	289	4.77
Business	10.30	–	–	357	69	288	2.24
Standard	2.44	–	–	1,167	1,166	1	7.31

Panel B: Consumer Surplus, Revenue, and Welfare (1,000 RMB)

	Consumer Surplus			Tax	Δ Road Cong.	Welfare		Δ Welfare (w/ MECC)	
	CS	ΔCS_ℓ	ΔCS_h	Rev.	Exter.	w/o MECC	w/ MECC	ΔW_ℓ	ΔW_h
Baseline	10,885	0	0	0	–	10,885	10,885	0	0
Crowding tax	7,671	-3,237	23	3,555	489	11,226	10,737	-707	559
Queuing	7,782	-3,008	-94	0	524	7,782	7,258	-3,469	-158
Two-class	11,737	27	825	0	-820	11,737	12,556	272	1,399

Notes: Panel A reports the consumer price P (equal to $P_{\text{base}} + T^*$ under the crowding tax), tax rate, queue wait W^* (in minutes), ridership totals by group, and crowding density k (persons/m²). Panel B reports consumer surplus and changes by group, tax revenue, the change in road congestion cost, and welfare with and without the MECC adjustment. The last two columns report group-level welfare changes inclusive of MECC:

$\Delta W_g = \Delta CS_g + \text{TaxRev}_g - \text{MECC}_g$, so $\Delta W_\ell + \Delta W_h$ equals the change in welfare with MECC relative to baseline. The queuing scenario targets the social optimum $N^* = 1,077$. Two-class pricing uses the variety-adjusted consumer surplus $CS_{\text{VA}} = CS_{\text{total}} - VB$. The revenue-neutrality target is $P_{\text{base}} \times N_{\text{eq}} = 6,533$ RMB. All counterfactuals use $\lambda = 0.5$ and the 3+3 car allocation.

generates 3,555 RMB in transferable revenue, so total welfare rises to 11,226 before accounting for road congestion. Once road-congestion costs from displaced riders are deducted, welfare is 10,737, slightly below baseline. The incidence is sharply regressive. Low-income consumer surplus falls by 3,237 RMB, while high-income consumer surplus actually rises by 23 RMB, because higher-income riders gain nearly as much from the reduction in crowding as they lose from the higher fare. Even when tax revenue is redistributed by ridership share, the net welfare change is -707 RMB for low-income riders and +559 for high-income riders.

Capacity-rationing queuing also targets N^* , with an equilibrium wait of $W^* = 7.75$ minutes, which is about 2.58 headways at the train frequency of 3 minutes. Consumer surplus is 7,782 RMB. Because the instrument generates no revenue, total welfare stays close to consumer surplus and far below the Pigouvian benchmark. The queuing deadweight loss, computed as $\sum_g N_g W^* (\text{VOT}_g + |\rho_g / \alpha_g|)$, is approximately 3,434 RMB, almost exactly the revenue that the Pigouvian tax would collect. This equivalence is the classical result from the queuing: rationing by time wastes as deadweight loss what rationing by price collects as transferable revenue.

Two-class pricing produces the highest welfare among all four policies. Revenue neutrality at $\lambda = 0.5$ implies optimal fares of $P_B^* = 10.30$ and $P_S^* = 2.44$ RMB. Total ridership rises to 1,524, an 18.5 percent increase over baseline, as the discounted Standard fare attracts new riders from the outside option. Variety-adjusted consumer surplus reaches 11,737 RMB, a 7.82 percent gain. The ridership expansion also reduces road congestion, adding 820 RMB in MECC savings. Variety-adjusted welfare including MECC is 12,556 RMB, 15.4 percent above baseline. In terms of distributional impacts, both groups gain on net once road-congestion benefits are redistributed by group: $\Delta W_\ell = 272$ and $\Delta W_h = 1,399$. Two-class pricing is the only instrument under which no group is worse off.

Three lessons emerge from the comparisons. First, transferable revenue is a first-order source of welfare for any policy that reduces ridership. The Pigouvian tax dominates queuing purely because its revenue can be redeployed, but queuing is slightly for progressive. Third, service differentiation combined with revenue-neutral pricing exploits the steep income gradient in crowding WTP documented in Section 6.3. High-income riders self-select into less crowded Business cars at a premium that cross-subsidizes lower Standard fares. Both groups gain on consumer surplus once road-congestion spillovers are included, so two-class pricing achieves a Pareto improvement.

8 Conclusions

This paper provides a revealed-preference estimate of subway passengers' willingness to pay to avoid crowding. Combining smartcard data from the Beijing Subway with an early-bird discount that creates plausibly exogenous variation in fares across departure times, we estimate that the marginal WTP to reduce in-train crowding by one passenger per square meter is about 0.053 RMB per minute. Aggregated over the average trip, this implies a crowding externality that exceeds the fare itself, suggesting that the unpriced congestion cost is a first-order distortion in subway pricing. The willingness to pay rises sharply with income while price sensitivity falls, so the average estimate masks substantial heterogeneity in how passengers value less crowded service. Passengers also dislike deviating from their preferred arrival time, with a baseline MWTP of 0.144 RMB per minute of schedule deviation. In the random coefficient model, this MWTP is about 0.10 RMB at the median income. Unlike crowding, this rescheduling cost is essentially flat across the income distribution, indicating that high- and low-income riders place similar value on arriving on time even as they differ markedly in how much they will pay to avoid a crowded car.

Building on these estimates, we evaluate three policies that target the social optimum. A Pigouvian crowding tax raises welfare relative to the unpriced equilibrium, but the gain is driven entirely by transferable revenue, and the incidence is sharply regressive. Capacity-rationing queuing produces consumer surplus close to the Pigouvian level but underperforms the tax because it generates no revenue, with the queue's wait cost almost exactly equal to the foregone fiscal transfer. A two-class configuration combined with revenue-neutral pricing dominates the alternatives. Higher-

income passengers self-select into a more expensive but less crowded Business car, the discounted Standard fare draws new riders from the outside option, and total welfare rises by roughly 15 percent above baseline once road-congestion spillovers are included. Both income groups are strictly better off, so two-class pricing is the only instrument under which no group bears a welfare loss.

Methodologically, the paper contributes to the literature on transit demand and congestion pricing along three dimensions. We use granular smartcard data to construct trip-level measures of in-train crowding from imputed routes, develop a novel instrumental variable based on overlapping trips between unrelated origin-destination pairs that addresses the endogeneity of crowding, and combine demand estimation with a calibrated welfare model that quantifies both within-system surplus and external road-congestion effects. From a policy standpoint, our results suggest that operators of crowded transit systems can achieve large efficiency gains and a Pareto improvement without raising average fares, provided that fare differentiation is accompanied by genuine differentiation in the riding experience.

References

- Anderson, Michael L.**, “Subways, Strikes, and Slowdowns: The Impacts of Public Transit on Traffic Congestion,” *American Economic Review*, 2014, 104 (9), 2763–2796.
- Arnott, Richard, Andre De Palma, and Robin Lindsey**, “A structural model of peak-period congestion: A traffic bottleneck with elastic demand,” *The American Economic Review*, 1993, 83 (1), 161–179.
- Beijing Municipal Institute of City Planning and Design**, “Trip-level Data in the Beijing Subway,” 2015.
- Beijing Transport Institute**, “Household Travel Surveys,” 2015.
- Berry, Steven, James Levinsohn, and Ariel Pakes**, “Automobile Prices in Market Equilibrium,” *Econometrica*, 1995, 63 (4), 841–890.
- Berry, Steven T.**, “Estimating discrete-choice models of product differentiation,” *The RAND Journal of Economics*, 1994, pp. 242–262.
- Buchanan, James M.**, “An Economic Theory of Clubs,” *Economica*, 1965, 32 (125), 1–14.
- Cook, Cody and Pearl Z. Li**, “Value Pricing or Lexus Lanes? The Distributional Effects of Dynamic Tolling,” *Working Paper*, 2025.
- Coulombel, Nicolas and Guillaume Monchambert**, “Diseconomies of scale and subsidies in urban public transportation,” *Journal of Public Economics*, 2023, 223, 104903.
- Davis, Lucas W.**, “Estimating the price elasticity of demand for subways: Evidence from Mexico,” *Regional Science and Urban Economics*, 2021, 87, 103651.
- de Palma, André, Robin Lindsey, and Guillaume Monchambert**, “The Economics of Crowding in Rail Transit,” *Journal of Urban Economics*, 2017, 101, 106–122.

- Douglas, Neil and George Karpouzis**, “Estimating the passenger cost of train overcrowding,” in “29th Australasian Transport Research Forum” 2006, pp. 1–17.
- Gordon, H. Scott**, “The Economic Theory of a Common-Property Resource: The Fishery,” *Journal of Political Economy*, 1954, 62 (2), 124–142.
- Gu, Yizhen, Qu Tang, Yacan Wang, and Ben Zou**, “Fare structure and the demand for public transit,” *Working Paper*, 2024.
- Hahn, Robert W, Robert D Metcalfe, and Eddy Tam**, “Welfare Estimates of Shifting Peak Travel,” Technical Report 31629, National Bureau of Economic Research 2023.
- Hall, Jonathan D**, “Inframarginal travelers and transportation policy,” *International Economic Review*, 2024, 65 (3), 1519–1550.
- Hardin, Garrett**, “The Tragedy of the Commons,” *Science*, 1968, 162 (3859), 1243–1248.
- Haywood, Luke and Martin Koning**, “The distribution of crowding costs in public transport: New evidence from Paris,” *Transportation Research Part A: Policy and Practice*, 2015, 77, 182–201.
- Hörcher, Daniel, Daniel J Graham, and Richard J Anderson**, “Crowding cost estimation with large scale smart card and vehicle location data,” *Transportation Research Part B: Methodological*, 2017, 95, 105–125.
- Hortaçsu, Ali, Olivia R Natan, Hayden Parsley, Timothy Schwieg, and Kevin R Williams**, “Organizational structure and pricing: Evidence from a large us airline,” *The Quarterly Journal of Economics*, 2024, 139 (2), 1149–1199.
- Huang, Ling and Martin D. Smith**, “The Dynamic Efficiency Costs of Common-Pool Resource Exploitation,” *American Economic Review*, 2014, 104 (12), 4071–4103.
- Kreindler, Gabriel**, “Peak-Hour Road Congestion Pricing: Experimental Evidence and Equilibrium Implications,” *Econometrica*, 2024, 92 (4), 1233–1268.
- Kumagai, Junya, Mihoko Wakamatsu, and Shunsuke Managi**, “Do commuters adapt to in-vehicle crowding on trains?,” *Transportation*, 2021, 48 (5), 2357–2399.
- Li, Shanjun**, “Better Lucky than Rich? Welfare Analysis of Automobile License Allocations in Beijing and Shanghai,” *Review of Economic Studies*, 2018, 85 (4), 2389–2428.
- Li, Zheng and David A Hensher**, “Crowding and public transport: A review of willingness to pay evidence and its relevance in project appraisal,” *Transport Policy*, 2011, 18 (6), 880–887.
- Lu, Hui, Tony Fowkes, and Mark Wardman**, “Amending the incentive for strategic bias in stated preference studies: case study in users’ valuation of rolling stock,” *Transportation research record*, 2008, 2049 (1), 128–135.
- Lu, Yao, Xinzheng Shi, Jagadeesh Sivadasan, and Zhufeng Xu**, “How Does Improvement in Commuting Affect Employees? Evidence from a Natural Experiment,” *The Review of Economics and Statistics*, 2024, 106 (1), 53–67.
- Parry, Ian W H and Kenneth A Small**, “Should urban transit subsidies be reduced?,” *American Economic Review*, 2009, 99 (3), 700–724.

- Ryan, Nicholas and Anant Sudarshan**, “Rationing the Commons,” *Journal of Political Economy*, 2022, 130 (1), 210–257.
- Singh, Jyotsna, Gonalo Homem de Almeida Correia, Bert van Wee, and Natalia Barbour**, “Change in departure time for a train trip to avoid crowding during the COVID-19 pandemic: A latent class study in the Netherlands,” *Transportation Research Part A: Policy and Practice*, 2023, 170, 103628.
- Small, Kenneth A. and Erik T. Verhoef**, *The Economics of Urban Transportation*, Routledge, 2007.
- Timmins, Christopher and Jennifer Murdock**, “A revealed preference approach to the measurement of congestion in travel cost models,” *Journal of Environmental Economics and management*, 2007, 53 (2), 230–249.
- Tirachini, Alejandro, David A. Hensher, and John M. Rose**, “Crowding in Public Transport Systems: Effects on Users, Operation and Implications for the Estimation of Demand,” *Transportation Research Part A: Policy and Practice*, 2013, 53, 36–52.
- Vickrey, William S**, “Congestion theory and transport investment,” *The American economic review*, 1969, 59 (2), 251–260.
- Wardman, Mark and Gerard Whelan**, “Twenty years of rail crowding valuation studies: evidence and lessons from British experience,” *Transport reviews*, 2011, 31 (3), 379–398.
- Weitzman, Martin L.**, “Prices vs. Quantities,” *Review of Economic Studies*, 1974, 41 (4), 477–491.
- Yang, Jun, Avralt-Od Purevjav, and Shanjun Li**, “The marginal cost of traffic congestion and road pricing: Evidence from a natural experiment in Beijing,” *American Economic Journal: Economic Policy*, 2020, 12 (1), 418–453.
- Yap, Menno, Oded Cats, and Bart Van Arem**, “Crowding valuation in urban tram and bus transportation based on smart card data,” *Transportmetrica A: Transport Science*, 2020, 16 (1), 23–42.

Crowding

Online Appendix (Not for Publication)

Yizhen Gu Qu Tang Kai Wu Ben Zou

Table of Contents

A	Data Sources	2
B	Market Size and Mode Choice Estimation	3
C	Additional Empirical Results	6
C.1	In-Train Crowding Density across Stations and Directions	6
C.2	Crowding and Travel Time	16
C.3	Evidence on Rerouting	19
D	Details of Welfare and Counterfactual Analysis	22
D.1	Numerical Implementation	22
D.2	Sensitivity to Car Allocation	23

A Data Sources

This appendix summarizes the data sources used in the paper.

Beijing Subway smartcard records. The primary data source is the trip-level records from the Beijing Subway’s automated fare collection (AFC) system (Beijing Municipal Institute of City Planning and Design, 2015). Each record captures the tap-in and tap-out stations and timestamps for a single trip, and trips can be linked across days through an anonymized card identifier. We estimate that between 90 and 95 percent of subway trips were paid by smartcard around the time of the EBD. The main sample consists of the universe of smartcard trips on 21 workdays in five non-consecutive collection periods in 2016, covering trips originating from the 16 EBD stations and the 16 control stations during the 6:30 to 8:30 AM window.

Subway network and operating information. We geocode all subway stations and collect the track distance and fare for each station pair from the Beijing Subway’s official website. We also collect train frequencies by line and station, the types of passenger cars, and the number of cars per train. These inputs are used to compute the total floor area available on each train segment at each moment and to determine the scheduled in-train travel time between adjacent stations. We treat seating and standing areas identically.

2015 Beijing Household Travel Survey. The Beijing Household Travel Survey (Beijing Transport Institute, 2015) is a large representative survey of approximately 100,000 households conducted by the Beijing Transport Institute. It records each respondent’s home and work locations, daily travel details, and transportation mode choices. We use the BHTS to fit the multinomial logit mode-choice model used in defining market sizes and the outside option.

Baidu Maps commuting flow data. We obtain grid-level commuting flow data from Baidu Maps, a leading digital map provider in China. The data track usual nighttime (home) and daytime (work) locations of roughly 7.6 million devices across about 15,000 residential grids with an edge of approximately 700 meters. For each home-work grid pair, the data report the share of commuters in each of five monthly income brackets. Income is imputed by Baidu’s internal machine-learning algorithm from device characteristics, installed applications, and location patterns. We use the Baidu data to construct OD-specific income distributions of potential subway riders.

Auxiliary sources. We collect supplementary information on the 16 EBD stations and their parallel control stations from the Beijing Subway’s public announcements. Official annual ridership statistics used to benchmark smartcard coverage are from the Beijing Municipal Commission of Transport.

B Market Size and Mode Choice Estimation

This appendix describes the multinomial logit mode-choice model fitted to the 2015 BHTS and the construction of OD-pair market sizes used in the demand estimation.

The Baidu commuting flow data capture trips in all modes of transportation. To define market sizes and the outside option in the demand model, we need to predict which commuters use the subway. We fit a multinomial logit model of mode choice using the 2015 BHTS, which includes a large representative sample of households with a rich set of demographic characteristics and mode choices for their trips.¹

We fit a multinomial logit where each observation is a commuting trip. The dependent variable is the main transportation mode, defined as the mode used for the longest leg. Modes include the subway, bus, and car, with other modes (bicycles, electric scooters, walking) as the other category. Explanatory variables include individual income, its square, and trip distance. Table B.1 reports the estimation results. Higher income is positively associated with all motorized modes relative to walking, biking, or scooting, with the largest effects for subway and car. Trip distance is also strongly positively associated with motorized modes.


Figure B.1 shows mode choice shares from the BHTS by income and trip distance during the morning peak (6 to 8 AM). The subway share rises sharply with trip distance, especially for middle- and high-income commuters. For commuters with an income above 8,000 RMB (the sample median), more than 40% choose subway for commutes longer than 12 km. In contrast, very few choose to ride the subway for trips under 5 km. Bus is primarily used by low-income individuals for middle-to-long distance trips. Car use is concentrated among higher-income commuters and increases with both income and distance, exceeding 30% for high-income commuters on longer trips. Subway nevertheless edges out for very long trips even at the highest income levels. Other modes, primarily walking and cycling, dominate only for short trips under 2 km and decline rapidly as distance grows, regardless of income.

To define market sizes, we use the fitted mode choice model to extrapolate to each OD pair. We predict user mode choices from the Baidu data using predicted mode shares based on income and travel distance. Let $\hat{s}_j^{\text{sub}}$ denote the subway mode share for OD pair j derived from the BHTS. The total market size is $M_{jd} = Q_{jd}^{\text{sub}} / \hat{s}_j^{\text{sub}}$, where Q_{jd}^{sub} is the observed subway ridership from the smartcard data. The outside good encompasses all non-subway modes and is given by $Q_{jd}^{\text{out}} = M_{jd} - Q_{jd}^{\text{sub}}$.

¹The BHTS reports only household income (in five bins), not individual income. We predict individual income by regressing household income (converted to log RMB) on flexible functions of each working household member's gender, age, education level, industry and occupation, their interactions, and a household fixed effect. The estimated coefficients are used to predict individual income subject to the constraint that all working members' predicted incomes sum to the observed household total.

Figure B.1: Mode Choice Shares by Income and Trip Distance

**Mode Choice Shares by Income and Trip Distance
(Beijing, Morning Peak 6-8 AM)**



Notes: Mode shares across income and trip distance bins from the BHTS 2015 survey (morning peak, 6 to 8 AM). Each cell shows the fraction of travelers in that income-distance bin choosing the given mode. Income is monthly household income per capita (1,000 RMB). Distance bins are based on quantiles of reported trip distance.

Table B.1: Estimation of Travel Mode Choice


	(1) Subway	(2) Bus	(3) Car
Income (1,000 RMB)	0.271 (0.007)	0.071 (0.004)	0.202 (0.004)
Income ²	−0.006 (0.000)	−0.001 (0.000)	−0.003 (0.000)
Trip Distance (km)	0.473 (0.004)	0.437 (0.004)	0.412 (0.004)
Observations	80,259		
Pseudo R ²	0.258		

Notes: Multinomial logit estimates. The baseline category is “other” modes (bicycles, electric scooters, walking). Standard errors in parentheses.

C Additional Empirical Results

C.1 In-Train Crowding Density across Stations and Directions

Figure C.1: In-Train Crowding Density



Notes: This figure depicts in-train crowding density k (persons/ m^2). The two graphs in the top row show downtown-bound trains over the course of the morning peak hours. Panel (a) averages across stations on EBD and control lines; panel (b) shows the most downstream station from each line. The two graphs in the bottom row show crowding density by station and direction for the two EBD lines during morning peak hours (6:30–8:30 AM). Stations are listed from the most upstream on the left to the most downstream on the right. Downtown-bound trains are substantially more crowded than suburban-bound trains. Crowding is computed using the full universe of rides.

Because stations in our sample sit at the far ends of suburban lines, the vast majority of morning passengers travel toward downtown. The top row of Figure C.1 show the average in-train crowding density during morning hours in downtown-bound trains passing through EBD and control stations. The two groups exhibit strikingly similar levels and time profiles. Crowding density averages about 1.5 persons per square meter around 7 AM, peaks at roughly 2.8 persons per square meter near 8:15 AM, and declines to 0.5 persons per square meter by 10 AM. The EBD and control series nearly overlay except in a narrow window around the 7 AM cutoff, where bunching at EBD stations creates a visible hump. The right sub-panel narrows the comparison to the most downstream station on each

line, where trains have accumulated passengers from all upstream stations. Although this sub-panel draws from only four stations (one from each line), the crowding patterns remain closely aligned between EBD and control lines. With passengers boarding from upstream stations, trains passing the most downstream stations are substantially more crowded. At around 8:20 AM, crowding density at the most downstream EBD stations peaks at approximately five persons per square meter.

The two graphs in the bottom row of Figure C.1 show that downtown-bound trains are far more crowded than suburban-bound trains on the two EBD lines. Trips in both directions qualify for the EBD and are included in our sample. The substantial difference in crowding between directions provides variation that helps identify potentially nonlinear crowding disutility at different density levels.


Passengers form expectations about crowding based on their commuting experience. To proxy for these expectations, we regress realized crowding on OD-by-day-of-week-by-time-bin fixed effects and use the fitted values $\widehat{\text{Crowd}}_{jt}$ as the expected crowding measure in the demand model. These fixed effects capture the systematic, predictable component of crowding patterns, precisely the information available to regular commuters. The same approach is applied to construct the expected travel time. The R^2 from this regression exceeds 0.98, indicating that passengers' crowding experiences are highly predictable from the OD pair, day of week, and time of day.

Table C.1: In-Station Time Estimation

	Estimated In-Station Time (min)				
	Mean	SD	P10	P90	N
Entry	3.50	1.84	1.36	5.34	33,436,600
Exit	2.00	0.90	0.89	3.15	33,436,600
Transfer	4.43	1.42	2.61	6.23	34,782,743


Notes: This table reports estimated in-station times from the iterative calibration procedure described in Section 4.1.1. Entry time is the duration from tapping in at the origin station to boarding the train. Exit time is the duration from alighting to tapping out, with the mean normalized to 2 minutes. Transfer time is the average duration across all transfer stations. N denotes the number of trips used in estimation.

Figure C.2: Beijing Subway Map with EBD and Control Stations




Notes: This figure shows the Beijing Subway network with the 16 early-bird discount (EBD) stations marked in red and the 16 control stations marked in blue. EBD stations are located on the Batong Line and the Changping Line. Control stations are on Line 5 and Line 6.

Figure C.3: Destination Distribution of EBD and Control Station Riders




Notes: Each point represents a destination station. The horizontal axis is the share of all control-origin riders going to that destination; the vertical axis is the corresponding share for EBD-origin riders. Blue dots denote the 188 destinations served by both groups; orange and green crosses denote the 8 EBD-only and 45 control-only destinations, respectively. The red dashed line is the 45-degree line. Pearson $r = 0.46$, Spearman $\rho = 0.68$ ($N = 241$).

Figure C.4: Simulated Monthly Income of Subway Riders (Thousand RMB)




Notes: Distribution of simulated monthly incomes (in thousands of RMB) drawn from OD-specific beta distributions fitted to the Baidu commuting data. The weighted mean income is 8.77 thousand RMB.

Figure C.5: Mean Monthly Income by Cell




Notes: Cell-level mean monthly income (thousand RMB, adjusted to 2016 levels) computed from the Baidu commuting portrait data as the population-weighted average across income brackets, with bracket midpoints as the per-bracket value. Cells are 2×2 aggregations of the underlying Baidu grid (cells of approximately 1.4 km on a side); cells with fewer than 5 commuters are dropped. The 2016 subway network (lines opened by 31 December 2016) is overlaid in red; the Beijing municipality boundary in dark grey.

Figure C.6: ICP Geography and Distance Distribution



Notes: Panel (a) shows the Beijing Subway network. EBD stations, located on the Changping Line (north) and the Batong Line (southeast), are highlighted in dashed boxes at the far suburban ends of the network. Red dashed circles mark 5-km, 8-km, and 10-km radii from the city center. The vast majority of the network lies beyond the 8-km ring, illustrating that the exclusion zone around a focal destination D removes only a small fraction of potential ICPs. Panel (b) shows the distribution of distances between each infrequent ICP's destination D' and the focal passenger's destination D . Dashed vertical lines mark the 25th, 50th, 75th, and 90th percentiles. The median ICP destination is approximately 11 km from D , well beyond the 8-km exclusion threshold used in the preferred IV specification. Infrequent riders are defined as those taking ≤ 3 trips in the current week.

Figure C.7: Increasing Marginal Disutility of Crowding



Notes: The figure uses the overall-sample specification with OD-by-date and date-by-time-bin fixed effects in Table 3, column (6). Utility includes squared expected cumulative crowding exposure, instrumented by expected ICT and the four EBD policy interactions. The horizontal axis divides cumulative exposure by the sample mean in-train duration to express it as an equivalent density. The solid line plots marginal WTP per person-minute per square meter; dashed lines show the 95% confidence interval.

Table C.2: Constant Coefficient Model Estimates: Frequent Riders

	OLS: Uninstrumented	GMM: ICT only		GMM: ICT + EBD instruments		GMM: Nonlinear
	(1)	1st Stage (2)	2nd Stage (3)	1st Stage (4)	2nd Stage (5)	2nd Stage (6)
Crowding	0.0062 (0.0003)		−0.0201 (0.0019)		−0.0094 (0.0011)	
Crowding ²						−0.000028 (0.000003)
Price	−0.2093 (0.0151)	−5.1519 (0.4331)	−0.3280 (0.0196)	29.3990 (2.1966)	−0.1821 (0.0094)	−0.1513 (0.0090)
Rescheduling	−0.0299 (0.0010)	0.0354 (0.0306)	−0.0274 (0.0014)	−0.1360 (0.0243)	−0.0250 (0.0010)	−0.0252 (0.0010)
ICT		11.5338 (0.4788)		11.1407 (0.4645)		
Other IVs				✓		
Wald <i>F</i> -Test			580.2		284.8	441.8
<i>R</i> ²	0.6588		−0.3272		−0.0720	
OD-Date FE	X	X	X	X	X	X
Date-Time FE	X	X	X	X	X	X
Mean Rescheduling (mins)	4.654		4.654		4.654	
Mean Crowding (person-min/m ²)	127.1		127.1		127.1	
<i>Implied MWTP</i>						
Crowding	−0.0296 (0.00266)		0.0613 (0.00550)		0.0516 (0.00613)	<i>see</i> <i>Fig. C.7</i>
Rescheduling	0.143 (0.0114)		0.0835 (0.00689)		0.137 (0.00894)	0.1662 (0.0120)

Notes: This table mirrors Table 3 in the main text but restricts the estimation sample to frequent riders, defined as passengers who take at least four subway trips in the current week. All specifications include OD-by-date and date-by-time-bin fixed effects. Column (1) reports uninstrumented OLS estimates. In columns (2)–(3), crowding is instrumented by ICT from infrequent riders with $|D - D'| \geq 8$ km. In columns (4)–(6), four EBD policy interactions are added as instruments (Treat $\times 1\{t < 7:00\}$, Treat $\times$ trend, Treat $\times$ trend², Treat $\times$ trend³). Column (6) includes a quadratic crowding term. Standard errors in parentheses, clustered at the OD level. MWTP is computed as the ratio of the relevant coefficient to the price coefficient. The estimates are within roughly ten percent of the overall-sample baseline reported in Table 3, confirming that the constant-coefficient results are stable across these two sample definitions. Crowding and ICT denote expected values; the row labels omit "expected".

Table C.3: Constant Coefficient Model Estimates Using ICP

	GMM: ICP only		GMM: ICP + EBD instruments	
	1st Stage (1)	2nd Stage (2)	1st Stage (3)	2nd Stage (4)
Crowding		−0.0175 (0.0019)		−0.0083 (0.0011)
Price	−4.9237 (0.4180)	−0.3075 (0.0185)		−0.1730 (0.0087)
Rescheduling	0.0305 (0.0279)	−0.0279 (0.0012)	−0.1365 (0.0218)	−0.0252 (0.0009)
ICP	178.1106 (8.0566)		179.4798 (7.8686)	
Other IVs			✓	
Wald F -Test		488.7		302.9
Observations	321,496	321,496	321,496	321,496
OD-Date FE	X	X	X	X
Date-Time FE	X	X	X	X
<i>Implied Willingness to Pay</i>				
MWTP for Crowding Reduction		0.0569 (0.0054)		0.0482 (0.0060)
MWTP for Resched. Avoidance		0.0907 (0.0071)		0.146 (0.0090)

Notes: This table reports GMM estimates using the number of incidental companion passengers (ICP) rather than incidental companion time (ICT) as the instrumental variable for crowding. Columns (1) and (3) report the first-stage regressions of in-train crowding on ICP under the ICP-only and ICP-plus-EBD specifications, respectively; columns (2) and (4) report the corresponding second-stage demand estimates with $\ln(\text{share})$ as the dependent variable. All columns include OD-by-date and date-by-time-bin fixed effects. The sample is the overall sample of riders. Standard errors in parentheses, clustered at the OD level. MWTP computed as the ratio of the relevant coefficient to the price coefficient. Crowding and ICP denote expected values; the row labels omit "expected".

Table C.4: Robustness: In-Station Crowding

	ICT only In-Station as Control (1)	ICT + EBD instruments In-Station as Endogenous (2)
In-train crowding	−0.0293 (0.0026)	−0.0078 (0.0013)
Price	−0.5855 (0.0323)	−0.1618 (0.0097)
Rescheduling	−0.0291 (0.0014)	−0.0242 (0.0011)
In-station crowding	0.0403 (0.0030)	−0.0077 (0.0023)
Mean in-train crowding (person·min/m ²)	127.09	127.09
Mean in-station crowding (thousand person·min)	14.31	14.31
Observations	295,400	295,400
Wald <i>F</i> -stat	424.40	470.58
MWTP in-train crowding (RMB per person·min/m ²)	0.0500 (0.0029)	0.0482 (0.0075)
MWTP in-station crowding (RMB per thousand person·min)	−0.0688	0.0476
OD-Date FE	X	X
Date-Time FE	X	X

Notes: Dependent variable is $\ln(S_{jdt})$ from the linear estimating equation (4), estimated by GMM. Standard errors clustered at the OD level. The in-station crowding regressor InStn_{jdt} is the product-level mean of passengers encountered during entry, transfer, and exit, weighted by the duration of overlap, expressed in thousands of person-minutes. Column (1) adds InStn_{jdt} as an uninstrumented control under the ICT-only instrument set; this specification is mis-specified because InStn_{jdt} is itself endogenous, and the positive coefficient on InStn_{jdt} reflects that endogeneity. Column (2) treats InStn_{jdt} as a second endogenous regressor, jointly instrumented with in-train crowding by the ICT-plus-EBD set: ICT (the infrequent-companion-time instrument, restricted to companions whose destinations lie at least 8 km from the focal D) plus the four EBD policy interactions $\text{treat} \times \mathbb{1}\{t < 7\text{am}\}$ and $\text{treat} \times t^k$ for $k = 1, 2, 3$. The sample consists of frequent riders. The two “Mean” rows report regression-sample averages of cumulative in-train crowding exposure (person·min/m²) and total in-station ridership-time (thousand person·min). WTP for in-train crowding reduction is $\hat{\beta}_{\text{crowd}}/\hat{\alpha}$, reported with delta-method standard errors. Crowding and ICT denote expected values; the row labels omit “expected”.

Table C.5: First Stage Estimation for Crowding (ICT + EBD Instruments)

	(1)	(2)	(3)
ICT	10.9955 (0.4318)	10.9076 (0.4441)	10.5954 (0.4185)
Treat $\times$ Pre-7am	51.7127 (3.4115)	51.7809 (3.4145)	35.3100 (3.4639)
Treat $\times$ Trend	534.2271 (34.8764)	533.2125 (34.8727)	590.7363 (38.4703)
Treat $\times$ Trend ²	-17.6784 (1.1699)	-17.6444 (1.1699)	-19.3772 (1.2872)
Treat $\times$ Trend ³	0.1930 (0.0130)	0.1926 (0.0130)	0.2096 (0.0143)
Price	29.1758 (2.0850)	29.2149 (2.0870)	19.1122 (2.0233)
Rescheduling	-0.1331 (0.0219)	-0.1363 (0.0221)	-0.1343 (0.0190)
Observations	321,496	321,496	321,496
OD-Date FE	X	X	X
Time FE	X		
Date-Time FE		X	
Time-Destin Line FE			X

Notes: This table reports the first stage of the ICT-plus-EBD GMM specification for the overall sample of riders. The dependent variable is expected in-train crowding (person·min/m²). Instruments are ICT from infrequent riders with $|D - D'| \geq 8$ km, plus four EBD policy interactions. Standard errors in parentheses, clustered at the OD level. Columns (1)–(3) differ only in the time fixed effects: time-bin (Column 1), date-by-time-bin (Column 2, used as the preferred specification), and time-bin-by-destination-line (Column 3). Crowding and ICT denote expected values; the row labels omit "expected".

C.2 Crowding and Travel Time

Column (1) of Table C.6 reports a regression of individual trip travel time on the average in-train crowding density (measured in persons per square meter) and the total number of in-station passengers (measured in counts), controlling for OD-by-date and line-time fixed effects. Each observation is an individual trip, and there are some 33 million observations. Increasing in-train crowding density by one person per square meter, an increase of roughly 35% from the mean, raises travel time by only 0.27 minutes. Given an average travel time of about 39 minutes, this effect represents a mere 0.7% increase. The regression includes the number of passengers the focal rider encountered during the trip, ensuring that the coefficient on crowding density captures the relationship between *in-train* crowding and travel time, which is the margin relevant for the demand model, rather than confounding it with delays caused by congestion in station corridors, platforms, and gates.

To put the magnitude in perspective, we convert the delay cost into monetary terms. The median monthly earnings for full-time workers is approximately 8,000 RMB. Assuming 160 work hours per month, this implies an earning rate of about 0.8 RMB per minute. Using the rule-of-thumb value of time in transit to be 50% of the wage rate, or 0.4 RMB per minute, the 0.27-minute increase in travel time from a one-person-per-square-meter increase in crowding translates to a monetary cost of about 0.11 RMB. As we show in the estimation results, this is an order of magnitude smaller than the willingness to pay for reducing discomfort from crowding.

The relationship between in-train crowding and delay is slightly convex. Column (2) of Table C.6 reports results from a quadratic specification. The linear term is negative (-0.16) and the quadratic term is positive (0.05), indicating that crowding has essentially no effect on travel time at low density levels and a small positive effect at high density levels. Figure C.8 plots the implied total monetary cost of crowding-induced delay as a function of crowding density. Below the 90th percentile of crowding density, the delay cost remains below 0.3 RMB, a small fraction of the average fare and the WTP for crowding reduction estimated later.

In-station crowding is also positively associated with total travel time, though it is not the focus of this paper. A one-person increase in in-station crowding is associated with a 0.003 minute increase in travel time. To compare the magnitudes, note that an increase of one person per square meter in in-train crowding density represents about a 35% increase from the ridership-weighted sample average of 2.87 persons/m². A comparable 35% increase from the mean in-station crowding of 668 persons corresponds to an increase of 234 persons. This translates to an increase in total travel time of $234 \times 0.0034 \approx 0.80$ minutes, or about 0.32 RMB in value of time. This is still small, but larger than that from in-train crowding. The fact that an increase in in-station crowding produces a larger delay than the same proportional increase in in-train crowding suggests that crowding delays travel primarily through slower walking traffic inside stations rather than by slowing trains down. This is intuitive, as pedestrian flows through corridors, stairways, and platform gates are directly impeded

Table C.6: The Effects of Crowding on Travel Time


	Travel Time	
	Linear (1)	Quadratic (2)
Crowding density k (persons/m ²)	0.2704 (0.0176)	-0.1586 (0.0468)
Crowding density k squared		0.0483 (0.0063)
In-station crowding (persons)	0.0034 (0.0000)	0.0037 (0.0001)
N	33,436,600	30,280,427
Mean dep. var.	38.59	38.05
Mean crowding density k (persons/m ²)	2.87	2.87
Mean in-station crowding (persons)	667.72	667.72
OD-Date FE	X	X
Line-Time FE	X	X

Notes: Regressions of trip travel time on in-train crowding density (persons/m²) and in-station crowding (number of passengers). Column (1) is the linear specification on the full sample. Column (2) is the quadratic specification, restricted to cells with at least two trips. Standard errors in parentheses, clustered at the OD level.

by larger crowds, whereas train speed is unaffected by passenger load.

While this paper focuses on in-train crowding, we note that in-station crowding also causes discomfort and generates disutility. In-station and in-train crowding are positively correlated, as a passenger experiencing a crowded platform can generally expect a crowded train. Although we do not directly estimate the disutility from in-station crowding, we show in the robustness checks that in-station crowding does not confound the estimated WTP to avoid crowded trains.

Figure C.8: Total Monetary Cost of Crowding-Induced Delay




Notes: This figure plots the implied delay cost (in RMB) from the quadratic specification in Column (2) of Table C.6 as a function of crowding density. Value of time is assumed to be 50% of median earnings (0.4 RMB/min). Delay cost remains below 0.3 RMB for crowding levels up to the 90th percentile.

C.3 Evidence on Rerouting

First, Figure C.9 shows the distribution of trip-level deviations from the OD-specific mean travel time. Panel (a) reports the deviation in minutes; panel (b) reports the ratio of individual travel time to the OD mean. The distributions are tightly concentrated around zero. The vast majority of trips cluster close to the mean, with 65.1% of trips falling within 3 minutes, 83.6% within 5 minutes, and 95.8% within 10 minutes. When expressed as a ratio to the mean, 71.6% of trips fall in the interval $[0.9, 1.1]$ and 94.8% in $[0.75, 1.25]$. If a substantial fraction of riders were taking alternative routes, we would expect to see a cluster of travelers in the far right tail, as rerouting causes considerable delay. No such cluster is apparent.

Figure C.9: Histograms of Deviation from the Mean Travel Time




Notes: Panel (a) shows the absolute deviation of each trip's travel time from the OD-specific mean (in minutes). Panel (b) shows the ratio of trip travel time to the OD-specific mean. Sample includes all trips in the working sample.

Alternative routes typically cost considerably more time than the shortest path, so if crowding systematically causes rerouting, it should have a large effect on the dispersion of travel time. Figure C.10 shows the distribution of travel time as a function of crowding. As in-train crowding increases, the distribution of trip time becomes slightly wider, indicating a larger variance, but the increase is not visually salient and there is no cluster of riders in the right tail at higher crowding quartiles. Much of the widening is driven by crowded stations slowing down foot traffic rather than by larger crowds slowing down trains. For trips that do not require a transfer, the distribution is much tighter and crowding does not lead to visually discernible increases in the variance. Trips requiring one or more transfers have larger variances in travel time to begin with, and crowding produces somewhat more visible increases in spread. Crucially, even for trips with transfers, the distributions remain similar across crowding quartiles, indicating that the right tail grows with the *number of transfers* rather than with *crowding levels*.

Second, we show formally in regressions that crowding has a negligible impact on the travel time variance. Table C.7 tests this directly. The dependent variables are the squared deviation (column

Figure C.10: Distribution of Travel Time Deviations by Crowding Quartile and Number of Transfers



Notes: Each row shows the distribution of deviations from the OD-specific mean travel time (in minutes) for a given transfer category, split by quartiles of in-train crowding density (least crowded on the left, most crowded on the right). Panel A includes all trips; Panel B includes trips without transfers; Panel C includes trips with one transfer; Panel D includes trips with two or more transfers. The distributions remain similar across crowding quartiles within each transfer category, indicating that the right tail grows with the number of transfers rather than with crowding levels.

1), absolute deviation (column 2), and ratio to mean (column 3) of each trip's travel time from its OD-date-time-bin cell mean. All regressions control for OD-by-date and line-time fixed effects and include both in-train crowding density and in-station ridership.

The coefficients on crowding density are statistically significant but economically small. In column 1, a one-person per square meter increase in crowding density raises the squared deviation of travel time from its cell mean by 0.60 minutes², about 3.7% of the sample mean of 16.38 minutes². The same one-unit increase is associated with a 0.12-minute increase in absolute deviation (column 2), relative to a mean absolute deviation of 2.40 minutes. In column 3, the correlation between crowding

Table C.7: The Effects of Crowding on Travel Time Variance

	(1) Squared Deviation	(2) Absolute Deviation	(3) Ratio to Mean
Crowding density k (persons/m ²)	0.6014 (0.1289)	0.1244 (0.0090)	-0.0000 (0.0000)
In-station ridership (persons)	0.0023 (0.0003)	0.0004 (0.0000)	0.0000 (0.0000)
N	30,280,427	30,280,427	30,280,427
Mean dep. var.	16.38	2.40	1.00
Mean crowding density k (persons/m ²)	2.87	2.87	2.87
Mean in-station ridership (persons)	667.72	667.72	667.72
OD-Date FE	X	X	X
Line-Time FE	X	X	X

Notes: Each observation is a trip. The dependent variable in column (1) is the squared deviation of the trip's travel time from its cell mean, $(TT_i - \bar{TT}_{\text{cell}})^2$, where a cell is defined as an OD-date-15-minute-bin. Column (2) uses the absolute deviation $|TT_i - \bar{TT}_{\text{cell}}|$. Column (3) uses the ratio of the trip's travel time to its cell mean, $TT_i / \bar{TT}_{\text{cell}}$. Cells with fewer than 2 trips are excluded, yielding 30,280,427 observations. Standard errors in parentheses, clustered at the OD level.

and ratio to the mean is essentially zero. These patterns are inconsistent with systematic rerouting, which would generate large right-tail outliers and substantial increases in travel time dispersion.

D Details of Welfare and Counterfactual Analysis

This appendix collects the numerical implementation details for the counterfactual exercises of Section 7 and reports sensitivity analyses with respect to the car allocation between Business and Standard classes.

D.1 Numerical Implementation

Our counterfactual computations rely on three numerical primitives: one-dimensional bisection for root finding, damped fixed-point iteration for the two-class ridership vector, and a scan-bracket-bisect procedure to enforce revenue neutrality. We summarize each here.

One-dimensional roots are computed by bisection. Given a continuous function F and an interval $[a, b]$ with $F(a)F(b) < 0$, we iterate $m = (a + b)/2$ and retain the sub-interval that preserves the sign change until $|F(m)|$ falls below a tolerance set between 10^{-9} and 10^{-12} depending on the application. We use this routine in three places: to invert the private demand mapping, to locate the Pigouvian social optimum, and to compute the queuing equilibrium wait.

The inverse private demand $PD(N)$ is evaluated as follows. For each candidate N , we bisect on P to solve $f(P; N) = \sum_g M_g s_g(A + \alpha_g P + \beta_g L N) - N = 0$. The function $f(\cdot; N)$ is continuous in P and strictly decreasing when $\alpha_g < 0$, which guarantees a unique root in a sufficiently wide bracket. The baseline calibration of A uses the same routine in reverse: we bisect on A so that the implicit solution of $PD(N) = P_{\text{base}}$ equals N_{eq} .

The two-class ridership vector (N_B, N_S) at prices (P_B, P_S) is obtained by damped fixed-point iteration. Starting from a feasible initial guess, we update

$$N_B^{(k+1)} \leftarrow (1 - \eta)N_B^{(k)} + \eta \sum_g M_g p_{gB}(V_{gB}(N_B^{(k)}), V_{gS}(N_S^{(k)})),$$

and analogously for $N_S^{(k+1)}$, with damping parameter $\eta \in [0.5, 0.7]$. Damping keeps the iterations stable when the nested-logit probabilities have steep local slopes in the class shares. Convergence is declared when both components change by less than 10^{-8} between successive iterations.

Revenue neutrality for the two-class counterfactual requires simultaneous control over P_B , P_S , and the equilibrium ridership vector. For each candidate P_S , we scan P_B upward from P_S to bracket every sign change in $R(P_B, P_S) - R_{\text{base}}$, where $R_{\text{base}} = P_{\text{base}} \times N_{\text{eq}} = 6,533$ RMB is the revenue target. Within each bracket we bisect to find a revenue-neutral P_B , and among all revenue-neutral pairs (P_B, P_S) we retain the one that maximizes the variety-adjusted consumer surplus. A final one-dimensional refinement over P_S yields the reported optimum.

Sensitivity to λ is reported in Table D.1. As λ rises, the variety bonus grows, and CS_{VA} falls toward baseline consumer surplus. At $\lambda = 0.8$, CS_{VA} is only +0.77 percent above baseline because revenue

neutrality forces $P_B \approx P_S$: the nested logit with a high λ penalizes unequal prices too strongly for self-selection to be sustained. At $\lambda = 0.3$, the gain reaches 15.6 percent, though such low nesting parameters are at the extreme end of the plausible range. The preferred value $\lambda = 0.5$ yields a 7.82 percent variety-adjusted welfare gain and is a conservative choice within the typical range.

Table D.1: Two-Class Pricing: Sensitivity to the Nesting Parameter λ

λ	P_B	P_S	N	VB	CS_{VA}	ΔCS_{VA} (%)
0.3	10.73	2.51	1538	1434	12579	15.56
0.4	10.63	2.48	1532	1918	12133	11.47
0.5	10.30	2.44	1524	2405	11737	7.82
0.6	9.74	2.43	1514	2896	11403	4.76
0.7	8.95	2.46	1505	3390	11144	2.38
0.8	7.76	2.64	1495	3888	10968	0.77

Notes: All scenarios maintain revenue neutrality ($R = 6,533$ RMB) with a 3+3 Business/Standard car allocation. VB denotes the variety bonus at the revenue-neutral equal price, $CS_{VA} = CS_{total} - VB$ is the variety-adjusted consumer surplus, and ΔCS_{VA} is the percentage change relative to baseline consumer surplus (10,885 RMB).

D.2 Sensitivity to Car Allocation

We examine how the two-class results vary with the Business/Standard split, holding $\lambda = 0.5$ fixed and maintaining revenue neutrality. A clear threshold emerges from Table D.2. When the number of Standard cars exceeds the number of Business cars, as in the 1+5 and 2+4 splits, the welfare optimizer fails to generate meaningful price differentiation. The revenue-neutral welfare-maximizing prices collapse to $P_B \approx P_S$, the price differential is numerically indistinguishable from zero, and the variety-adjusted consumer surplus gain vanishes. The mechanism is direct. When the Standard class has excess capacity, per-rider crowding in Standard is already low at the revenue-neutral equal price, so riders have little incentive to pay a premium for Business class. The optimizer cannot sustain self-selection because the crowding differential between classes disappears mechanically.

Splits with $J_B \geq J_S$ produce substantial price differentiation. Under the current calibration, the preferred 3+3 allocation yields $P_B^* = 10.30$, $P_S^* = 2.44$, with a variety-adjusted consumer surplus gain of 7.82 percent (Table D.1, row $\lambda = 0.5$). Other splits in Table D.2 (1+5, 2+4, 4+2, 5+1) have not been re-run under the current calibration; the qualitative threshold remains: revenue-neutral two-class pricing is effective only when $J_B \geq J_S$. Below this threshold, revenue neutrality forces the two-class system to collapse to a single uniform fare.

Table D.2: Two-Class Pricing: Sensitivity to Car Allocation

	$J_S > J_B$		3+3	$J_B \geq J_S$	
	1+5	2+4		4+2	5+1
P_B^* (yuan)	4.46	4.43	6.98	5.77	4.80
P_S^* (yuan)	4.43	4.43	2.97	2.96	3.57
$P_B^* - P_S^*$ (yuan)	0.03	0.00	4.01	2.81	1.22
Total ridership (N)	1,506	1,508	1,543	1,560	1,523
Price discrim. gain (RMB)	-2.7	-0.6	108.1	231.1	58.0
ΔCS_{VA} (%)	-0.03	-0.01	+1.17	+2.50	+0.63

Notes: All scenarios maintain revenue neutrality ($R = 6,533$ RMB) with $\lambda = 0.5$ and $J_B + J_S = 6$. Price discrimination gain is the difference between CS at the welfare-maximizing revenue-neutral prices and CS at the revenue-neutral equal price. ΔCS_{VA} is the percentage change relative to baseline single-class consumer surplus (10,885 RMB; the value shown for the 3+3 row reflects the current calibration, while the other rows are pre-recalibration and have not been regenerated).